

Week 6: Avoiding Locks & Scalability

Non-blocking Synchronization

Read-Copy Update

Transactional Memory

Scalability

CSC 469 / CSC 2208

Fall 2018

(with thanks to Bogdan Simion, Tom Hart, Paul McKenney)



University of Toronto, Department of Computer Science



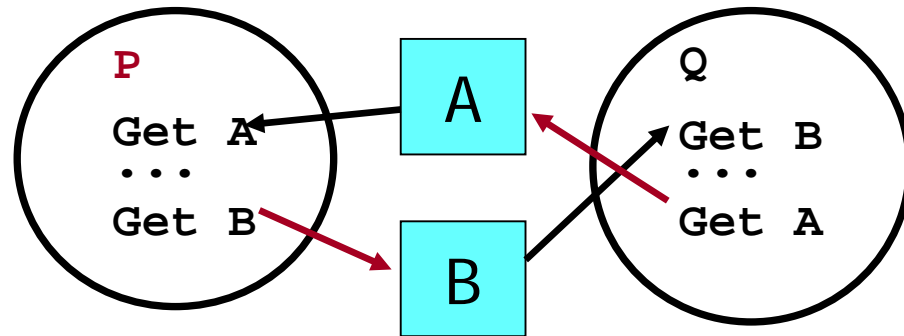
Locking: A necessary evil?

- Locks are an easy to understand solution to critical section problem
 - Protect shared data from corruption due to simultaneous updates
 - Protect against inconsistent views of intermediate states
- But locks have lots of problems
 - 1. Deadlock
 - 2. Priority inversion
 - 3. Not fault tolerant
 - 4. Convoying
 - 5. Expensive, even when uncontended
- *Not* easy to use correctly!



1. Deadlock

- Textbook definition: Set of threads blocked waiting for event that can only be caused by another thread in the same set
- Classic example:




- Self-deadlock also a big issue
 - Thread holds lock on shared data structure and is interrupted
 - Interrupt handler needs same lock!
- Solutions exist (e.g., specify lock order, disable interrupts while holding lock) but add complexity



2. Priority Inversion

- Lower priority thread gets spinlock
- Higher priority thread becomes runnable and preempts it
 - Needs lock, starts spinning
 - Lock holder can't run and release lock



- Solutions exist (e.g. disable preemption while holding spinlock, implement priority inheritance, etc.), but add complexity



3. Not fault tolerant

- If lock holder crashes, or gets delayed, no one makes progress

✓ lock
~~CRASH!~~

✗ lock
spin


✗ lock
spin

- Scheduler-conscious synchronization helps with delays (preemption, page faults)
 - Crashes require abort / restart



4. Convoying

- Threads doing similar work, started at different times, occasionally accessing shared data
 - e.g., multi-threaded web server
 - Expect access to shared objects to be spread out over time
 - Lock contention should be low
 - Delay of lock holder allows other threads to catch up
 - Lock becomes contended and tends to stay that way
- => **Convoying**





5. Expensive, even when uncontended

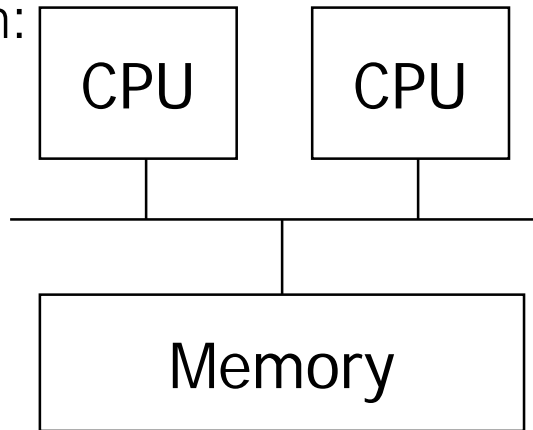
| Operation | Nanoseconds |
|---------------------------------------|-------------|
| Instruction | 0.24 |
| Clock Cycle | 0.69 |
| Atomic Increment | 42.09 |
| Cmpxchg Blind Cache Transfer | 56.80 |
| Cmpxchg Cache Transfer and Invalidate | 59.10 |
| SMP Memory Barrier (eieio) | 75.53 |
| Full Memory Barrier (sync) | 92.16 |
| CPU-Local Lock | 243.10 |

McKenney, 2005 – 8-CPU 1.45 GHz PPC

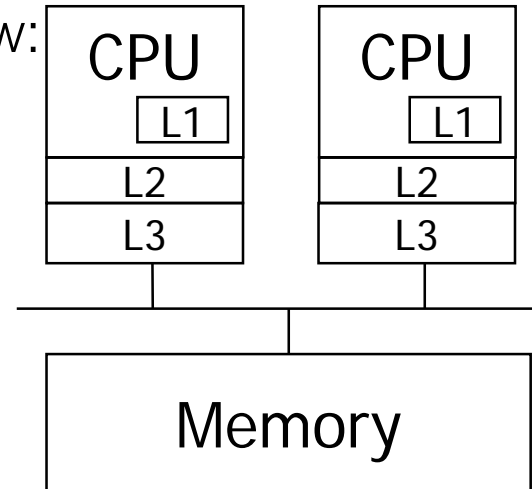


Causes: Deeper Memory Hierarchy

Then:



Now:



- Memory speeds have not kept up with CPU speeds
 - 1984: no caches needed, since instructions slower than memory accesses
 - after ~2005: 3-4 level cache hierarchies, since instructions orders of magnitude faster than memory accesses
- Synchronization ops typically execute at memory speed



Causes: Deeper Pipelines

Then:



Now:



- 1984: Many cycles per instruction
- 2005: Many instructions per cycle
 - 20 stage pipelines
 - CPU logic executes instructions out-of-order to keep pipeline full
 - Synchronization instructions must not be reordered
 - => synchronization stalls the pipeline
- Deeper pipelines not always better and processors are changing




Performance

- Main issue with lock performance used to be contention
 - Techniques were developed to reduce overheads in contended case
 - And to reduce contention
- Today, issue is degraded performance even when locks are *a/ways* available
 - Together with other concerns about locks



Critical section efficiency

- Assuming little to no contention, and no caching effects in CS



$$Efficiency = \frac{T_c}{T_c + T_a + T_r}$$

- Even if lock contention is negligible, critical section efficiency must be addressed!



Locks: ~~A necessary~~ evil?

Idea: Don't lock if we don't need to!

- Non-Blocking Synchronization (NBS)
 - Use term “*lockless*” to describe strategies that avoid locking



NBS Basics

- Make change optimistically, roll back and retry if conflict detected

```
atomic_inc(int *counter) {  
    int value;  
    do {  
        value = *counter;  
    } while (!CAS(counter, value, value+1));  
}
```

- Complex updates (e.g. modifying multiple values in a structure) are hidden behind a single commit point using atomic instructions



Example: Stack Data Structure

- Lock-based synchronization:

```
/* definitions */

typedef struct node_s
{
    int val;
    struct node_s *next;
} node_t;

typedef struct stack_s
{
    node_t *top;
    lock_t *stack_lock;
} stack_t;
```

```
void push(stack_t *S, node_t *n)
{
    lock(S->stack_lock);
    n->next = S->top; S->top=n;
    unlock(S->stack_lock);
}

node_t* pop(stack_t *S) {
    node_t *n = NULL;
    lock(S->stack_lock);
    if (S->top != NULL) {
        n = S->top;
        S->top = S->top->next;
    }
    unlock(S->stack_lock);
    return n;
}
```



Non-blocking stack (take 1)

```
/* definitions */

typedef struct node_s {
    int val;
    struct node_s *next;
} node_t;

/* Stack type is just
 * a pointer to a node.
 */
typedef
    node_t* stack_t;
```

What's wrong?



```
void push(stack_t *S, node_t *n)
{
    node_t *first;
    do {
        first = *S;
        n->next = first;
    } while (!CAS(S,first,n));
}

node_t* pop(stack_t *S) {
    node_t *first, *second;
    do {
        first = *S;
        if (first != NULL) {
            second = first->next;
        } else return NULL;
    } while (!CAS(S,first,second));
    return first;
}
```



ABA Problem

- T_i, T_j both doing pops and pushes, interleaved as follows:

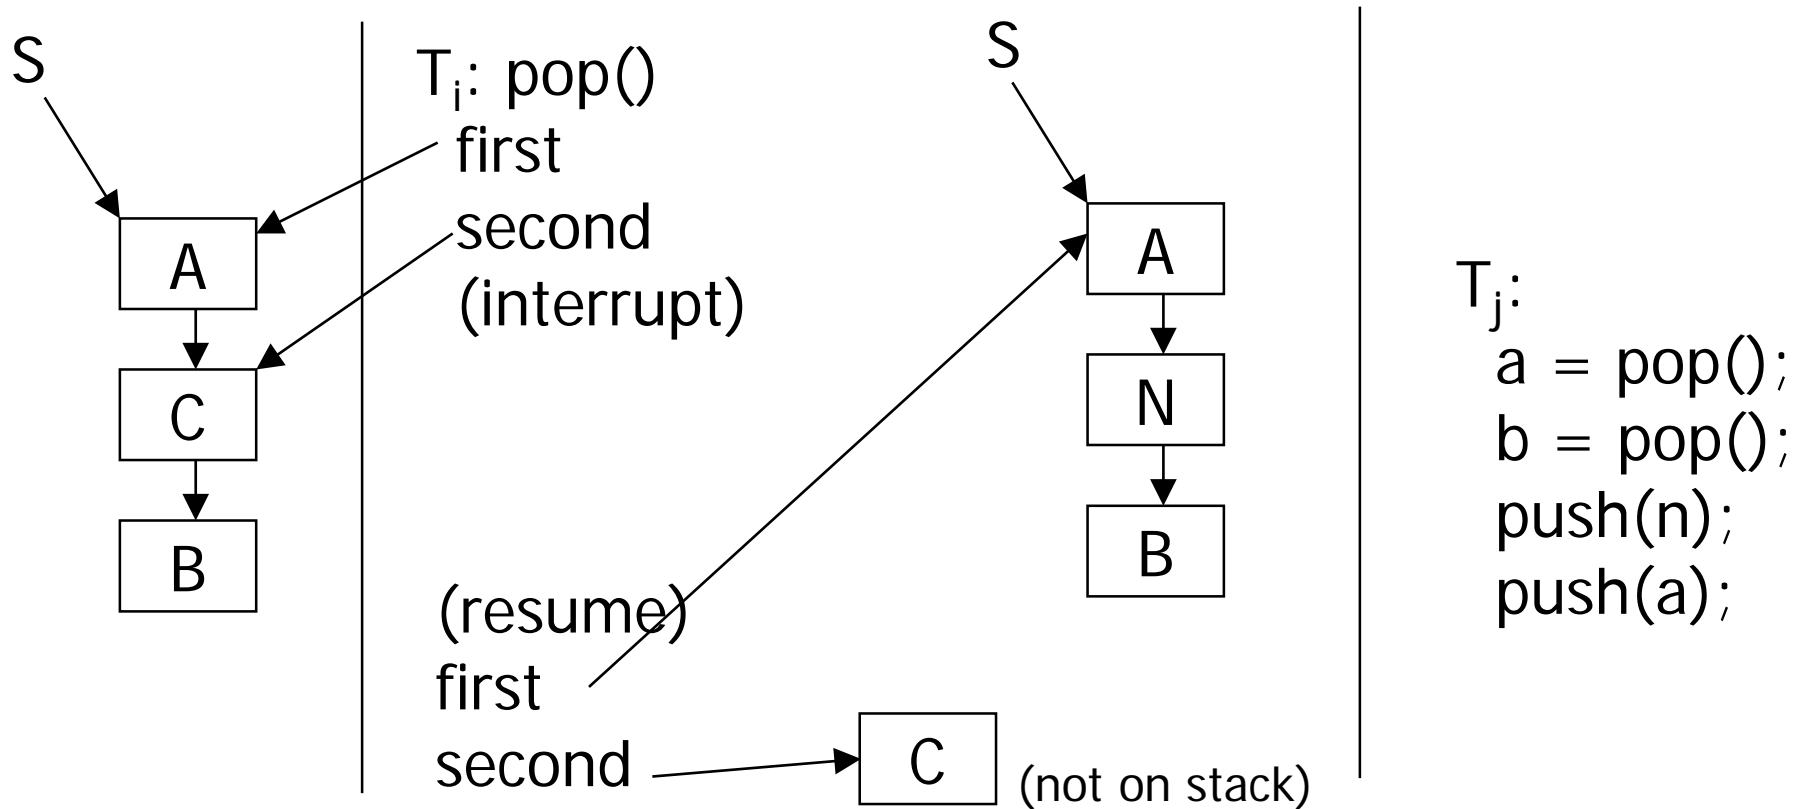


T_j :
a = pop();
b = pop();



ABA Problem

- CAS(x, y, z) succeeds if value stored at x matches y





One Solution

- Include a version number with every pointer
 - `pointer_t = <pointer, version>`
 - Increment version number (atomically) every time you modify pointer
 - Change to version number guarantees CAS will fail if pointer has changed
 - Requires double-word CAS operation (most architectures do not provide this)
 - Use garbage collection to reclaim memory later
 - May restrict reuse of memory



Using NBS

- Good for simple data structures, update heavy
- When you need NBS constraints/guarantees
 - Progress in face of failure
 - Linearizability
 - Everyone agrees on all intermediate states
- Both constraints are often irrelevant!



Constraints Irrelevant?

- Real systems don't fail the way theoretical ones do
 - Software bugs are not always fail-stop
 - Preemption/interrupt is not a failure
 - And can be controlled by system programmer or scheduler-conscious synchronization
 - Page fault is not a failure
 - Over-provision memory... if shared data really is paged out, it will have to be brought into memory before progress is made anyway
- Don't always need intermediate states, just final
 - Linearizability implies dependency → limits parallelism
 - If events are unrelated, asynchronous, does it matter which happened first?



Read-Copy Update (RCU)

- What is RCU?
 - Paul McKenney's PhD thesis
 - a key part of the Linux scalability effort
 - and one of the key technologies in the SCO lawsuit against IBM.
- Ok, what is it really?
 - Reader-writer synchronization mechanism
 - Readers use no locks; best for read-mostly data structures
 - Writers create new versions atomically
 - typically by locking out other writers
 - Readers can continue to access old versions
 - Old versions must be deleted at some point
 - “poor man's garbage collection”



RCU Basics

- From <http://lwn.net/Articles/262464>
 1. Publish/Subscribe mechanism (for insertion)
 2. Mechanism to wait for previous readers to complete (for deletion)
 3. Maintain multiple versions of recently updated objects (for readers)



RCU Publish/Subscribe

```
/* definitions */
struct foo {
    int a;
    int b;
    int c;
};

struct foo *gp = NULL;
```


```
T1: p = kmalloc(sizeof(*p),
                GFP_KERNEL);
T1: p->a = 1;
T1: p->b = 2;
T1: p->c = 3;
T1: gp = p; rcu_assign_pointer(gp,p);
...
    rcu_read_lock();
T2: p = gp; p = rcu_dereference(gp);
T2: if (p != NULL)
T2:     use(p->a, p->b, p->c);
    rcu_read_unlock();
```

- When is it safe to read a pointer?
 - RCU Readers use no locks
 - Compiler, CPU may reorder assignments
 - Enforce ordering with rcu_assign_pointer/rcu_dereference



RCU Deletion Example

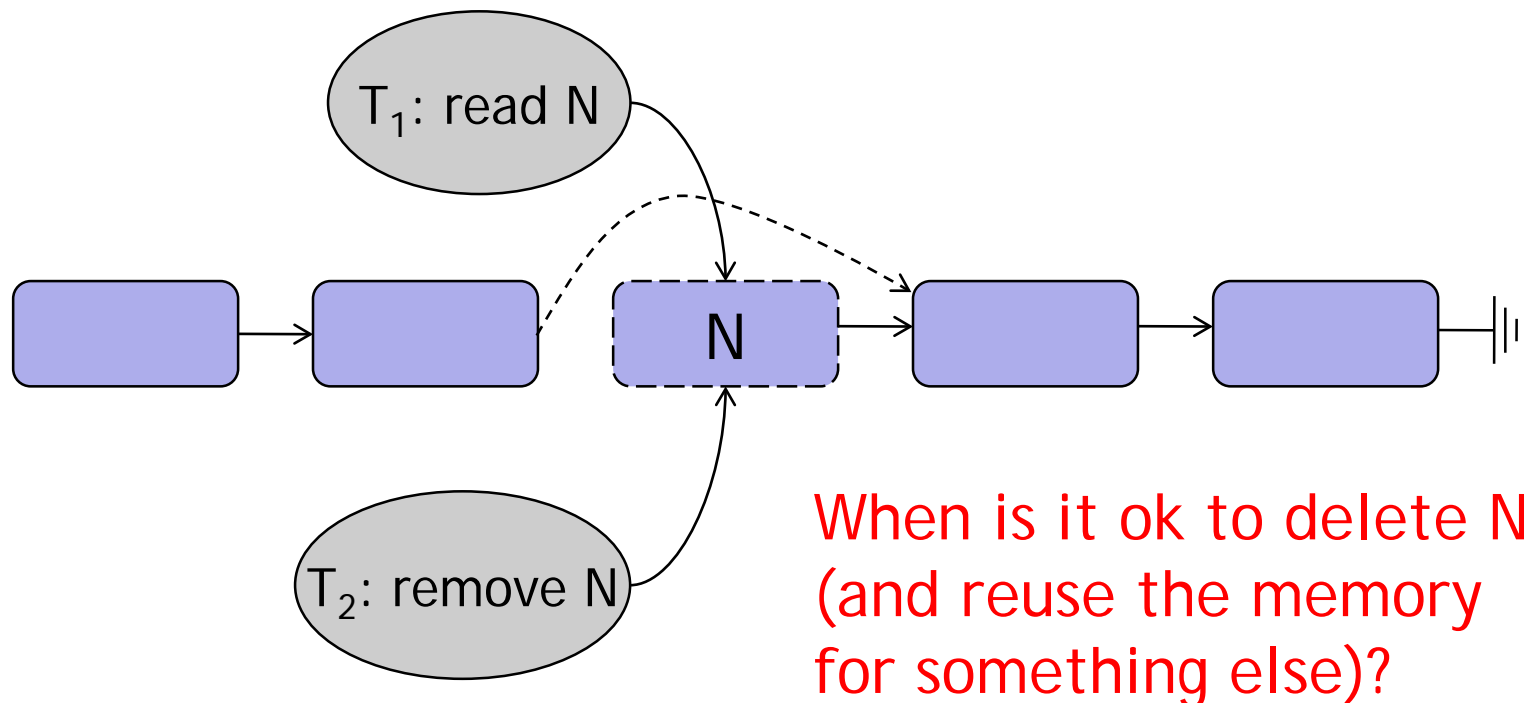
- T_1 traversing linked list, T_2 removes an element:





RCU Deletion Example (2)

- After removal – T_1 continues to use N and later nodes in the list



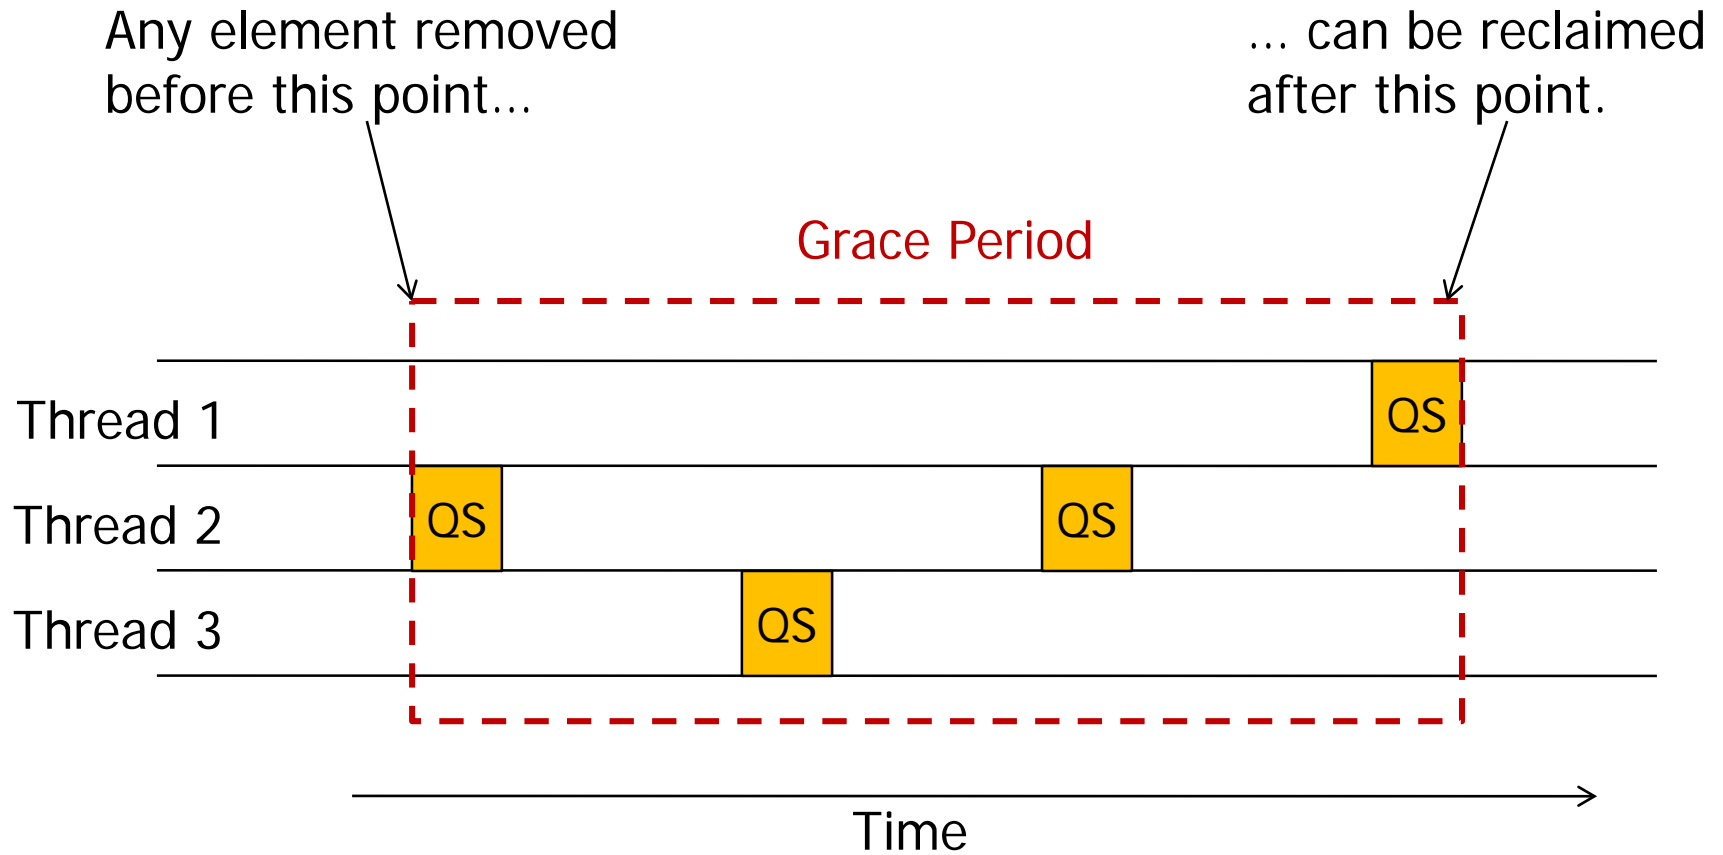


Handling read-reclaim races

- RCU uses *quiescent state based reclamation* (QSBR)
- **Defn:** A *quiescent state* for a thread T is a state in which T holds no references to shared data
- **Defn:** A *grace period* is an interval in which every thread has passed through at least one quiescent state
- **Basic Idea:** elements removed from a data structure can be reclaimed after a grace period, since no thread can still be holding a reference to the old element at that point



Illustration






How to define Quiescent States?

- Application dependent!
- For OS kernels, some natural ones exist
 - E.g. a context switch in a non-preemptive kernel
- RCU primitives
 - `rcu_read_lock()` and `rcu_read_unlock()`
 - Surround read-side critical sections
 - No overhead (`#define'd` as nothing) in non-preemptive kernels
 - Modest overhead in preemptive kernels (disable preemption)
 - `synchronize_rcu()`
 - Wait until all pre-existing RCU read-side critical sections complete
 - Force execution on all CPUs



PPC Hash Table with RCU





When to use which tool

- Read-mostly situations
 - RCU (if algorithm can tolerate concurrent reads and updates)
- Update-heavy situations
 - Simple data structures and algorithms: NBS
 - Complex data structures and algorithms: Locking

“When the only tool you have is a hammer, everything looks like a nail.”

- It's good to have lots of tools in your toolbox



Transactional Memory

Active research! Here be dragons...





Challenges of Synchronization

- Two major issues:
 - Performance
 - Scalability
 - Base cost
 - We have looked at some techniques that address this
 - Better spinlocks
 - Lockless strategies (NBS, RCU)
 - Programmability
 - Locks are hard to use correctly
 - Lockless data structures are hard to design



What's missing?

- Lack of support for *abstraction* and *composition*
- E.g. Suppose we have thread-safe stack with (abstract) push and pop operations
 - In sequential programs, can use these operations without regard to their implementation
 - In parallel programs, internal details may be needed
 - Consider task of moving an item from one stack to another
 - Need to expose stack locking mechanism



“Magic” Wish List

- Let programmers express desired outcome
 - “This block of code should appear atomic”
- Let run-time system or hardware support make it happen
- Allow abstractions to hide implementation and be composable



A new programming model is needed



Database Transactions

- Database systems allow multiple queries to run in parallel
- Query authors don't worry about concurrency
- Complex queries can be composed out of simpler ones
- Can we use the DB programming model as a general parallel model?
- **Key Programming Model:** everything is a transaction
 - A transaction executes as if it were the only computation accessing the database
 - Restricted interactions, serializability
 - Hide complex implementation detail, programmer only sees a simple interface
 - **A**tomic – all updates become visible, or none
 - **C**onsistent – transactions leave database in consistent state
 - **I**solated – no interference with or from other transactions
 - **D**urable – once committed, updates are permanent



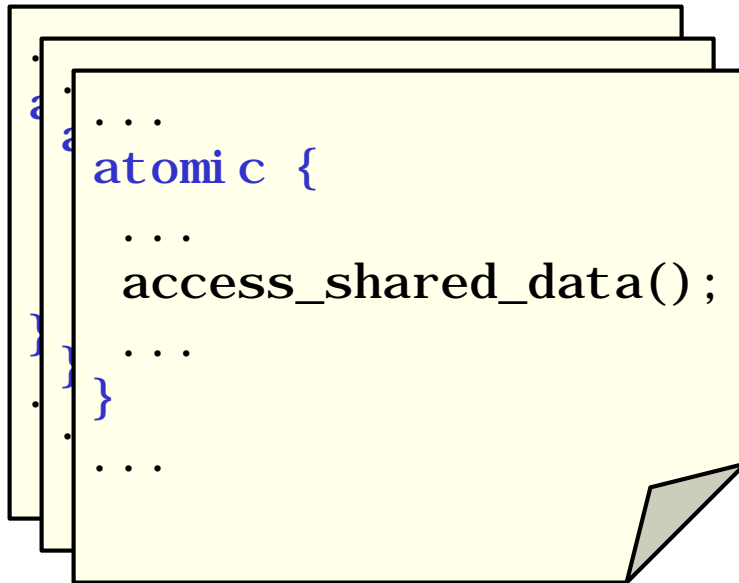
Transactional Memory: Some History

- 1977 – D.B. Lomet (IBM Research, now at Microsoft Research) suggests database transaction model for concurrent programming
 - No practical implementation provided
- 1983 – Kung & Robinson propose *optimistic concurrency control* for databases
- 1988 – Chang & Mergen describe IBM 801 storage manager
 - HW provided lock bits for each 128 byte range of a page; page tables & TLB extended
- 1993 – Herlihy & Moss describe a hardware proposal for *transactional memory*




Transactional Memory (TM)

Source Code:



TM System

Transactions:



Programmer: Specifies threads/transactions in source code

TM System: Executes transactions optimistically in parallel

- 1) Checkpoints execution
- 2) Detects conflicts
- 3) Commits or aborts and re-executes



Differences from DB Transactions

- Memory vs. disk
 - Disk access takes 100X longer than memory access
 - database systems can use relatively heavy-weight software solutions
- No need for durability
 - Memory is transient anyway
 - simplifies TM implementations
- Existing languages, libraries and systems
 - Databases are closed systems in which all code executes as a transaction, BUT programs using TM must coexist with libraries and OSs that do not



TM Implementations

- Hardware TM (HTM)
 - Changes to computer system and ISA
 - Extra cache to buffer writes, extended coherence protocol to track conflicts, special transaction instructions
 - Support for limited number of memory locations
- Software TM (STM)
 - Language runtime (or library) + extensions to specify transaction
 - Exploit current commodity hardware (multicores)
 - Get experience with transactional programming model
 - Java: DSTM (Marathe et al.), ASTM (Herlihy et al.)
 - C/C++: McRT-STM (Saha et al.), TL2 (Dice et al.), RSTM
 - Intel's C++ STM compiler
- Hybrid TM (HyTM)



Programming Constructs

- Atomic block

```
atomic {  
    if (x!=null) x.foo();  
    y = true;  
}
```

- Delimits code that should execute in a transaction
- Dynamically-scoped – code in foo() executes in transaction as well
- Does not name shared resources (unlike monitors or lock-based programming)
- 3 possible outcomes – commits, aborts, non-termination



Caution!

- Programmers can still use *atomic* incorrectly

```
bool flagA=false; bool flagB=false;
```

Thread 1:

```
atomic {  
    while (!flagA);  
    flagB = true;  
}
```

Thread 2:

```
atomic {  
    flagA = true;  
    while (!flagB);  
}
```

- What's wrong?
 - Deadlock results



Semantics

- Not yet formally specified!
- Useful ways to reason about TM:
 - Database correctness criteria: serializability
 - Useful for understanding transaction behaviour
 - Says nothing about interaction of transactions with code outside of transactions
 - Operational semantics – single-lock atomicity (SLA)
 - Program executes as if all atomic blocks were protected by single global lock
 - Attractive, but may be problematic conceptually
 - SLA does not support failure atomicity, forms of nesting, etc.



Implementation Basics

- For all (non-stack) write instructions:
 - Track write addresses and values (*write set*)
- For all (non-stack) read instructions:
 - track read addresses and values (*read set*)
- When a transaction completes:
 - Atomically
 - Validate read set (conflict detection)
 - Commit write set



Implementation Options

- Transaction Granularity
 - Unit of storage over which TM system detects conflicts
 - Similar to notion of cache coherence
 - Word or block typical for HTM, object common for STMs that extend OO language
- Direct or Deferred Update
 - Direct – transaction directly modifies the object itself
 - Must log previous value for undo in case of abort
 - Deferred – modify private copy, propagate at commit
 - Both get complicated in the presence of data races
- Optimistic or Pessimistic Concurrency Control
 - TM typically optimistic; need to detect and resolve conflict




Location-Based Conflict Detection

Transaction 1:



Strip versions:

Main Memory:



Strip versions:

Transaction 2:



Strip versions:



Strips

Legend:



Read




Written



Location-Based Conflict Detection


Transaction 1:

Strip versions:



Main Memory:

Strip versions:



Transaction 2:

Strip versions:



Legend:



Read




Written



Location-Based Conflict Detection


Transaction 1:

Strip versions:




Main Memory:

Strip versions:



Transaction 2:

Strip versions:



Commit step 1) Validate Read Set ✓

Commit step 2) Publish Writes (and inc version #s)

Legend:



Read



Written



Location-Based Conflict Detection


Transaction 1:

Strip versions:




Main Memory:

Strip versions:



COMMITTED
Transaction 2:

Strip versions:



Commit step 1) Validate Read Set **X** Abort!

Note: all transactions must maintain strip version #s

Legend:



Read




Written




Value-Based Conflict Detection

Transaction 1:



Main Memory:



Transaction 2:



Legend:



Read




Written




Value-Based Conflict Detection

Transaction 1:



Main Memory:



Transaction 2:



Legend:



Read




Written




Value-Based Conflict Detection

Transaction 1:



Main Memory:



Transaction 2:



Commit step 1) Validate Read Set ✓

Commit step 2) Publish Writes

Legend:



Read




Written




Value-Based Conflict Detection

Transaction 1:



Main Memory:



Transaction 2:



Commit step 1) Validate Read Set **X** Abort!

Note: no version information to maintain

Legend:



Read



Written



TM Weaknesses

- Some operations are hard to abort/retry
 - Essentially anything not idempotent, e.g. I/O
- In practice, TM does not interact well with locking
- Some variables are prone to high conflict rates (frequent true sharing & dependences)
- Conflict resolution needs to avoid starving long-running, large transactions
- Poor interaction with standard software tools like debuggers
 - Getting better though ...



TM Status

- Hardware TM is now a reality
 - Sun's Rock processor was killed after acquisition by Oracle (2009)
 - Azul Systems has HTM in their Java appliance hardware (circa 2009)
 - IBM BlueGene/Q (2011)
 - Intel Haswell's *Transactional Sync Extensions (TSX)*
- Software TM has performance problems
 - But some applications are a nice fit
 - E.g. parallel game server




Operating System Scalability

- We have looked at various synchronization strategies
 - Scalability has been a key concern
- Most user applications actually aren't very scalable
- Most exceptions use few OS services anyway
 - E.g. scientific computing
- Most multiprocessor systems support independent processes (multi-user workloads)
- Why does OS scalability matter?




Systems View of Scalability

User Commands




Scale up




Add
Scripts &
Processors
(SMP)

Users Commands



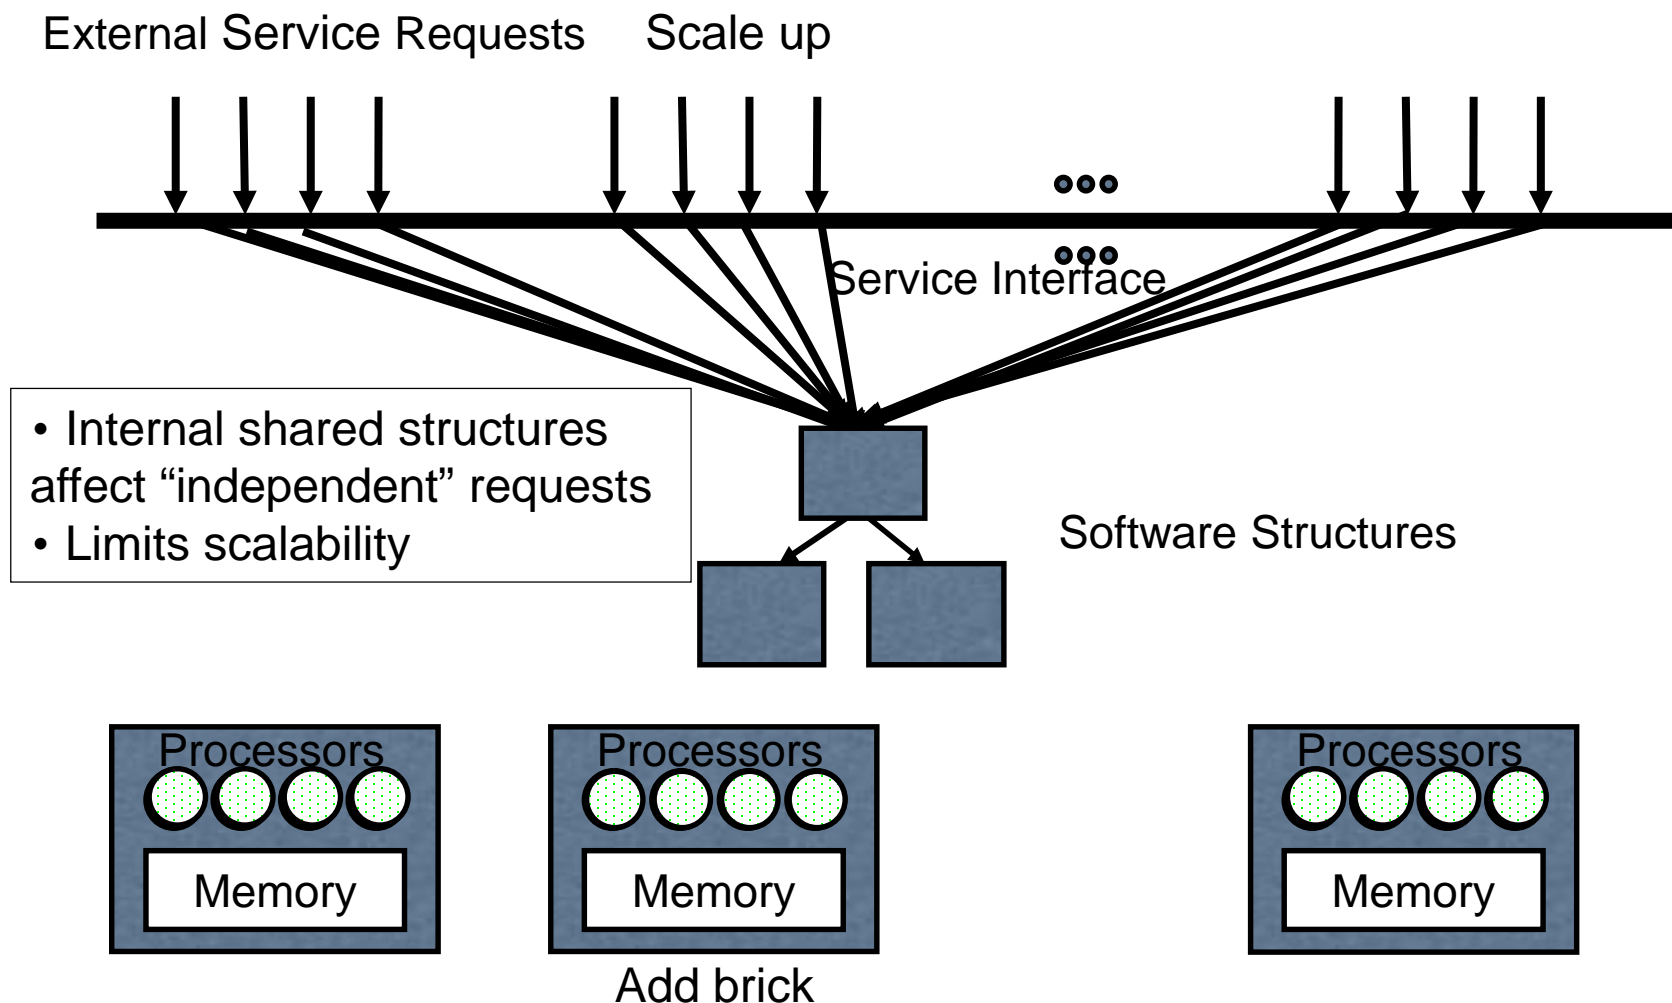


The Problem





Scaling Existing OSes






Areas of Concern

- Statistical counters
 - Widely used to track variety of system properties
 - Frequently updated, rarely read
- Processor scheduling
 - We'll look at this closely in the next 2 lectures
- Memory management
 - In tutorial this week (also, Assignment 2!)




Simple Shared Counter Example






Solution: Per-CPU data

- OS assigns each CPU an integer *id* at boot time
 - Linux: access with `smp_processor_id()`
- Basic data structure is array with entry for each CPU
 - `counter[smp_processor_id()]` is data structure for current CPU





Simple Shared Counter Example






What went wrong?

- Per-CPU array can lead to *false sharing* problem
 - Each CPU has own variable
 - Several per-CPU variables are on same cache line
 - Modification of one causes invalidates in other CPUs' caches
- Solutions?
 - Use *padding* so each per-CPU variable lies on different cache line



Simple Shared Counter Example





Summary

- Taking a traditional OS and making it scale well on shared memory multiprocessors is hard
 - Fast uniprocessor solutions typically don't scale
 - Designing for scalability can hurt uniprocessor performance
 - Maintaining scalability with every change is hard
- => Must design a system from the ground up, with scalability in mind



Insights and Approaches

- Scalability must be considered in system design
- Shared data is the enemy
 - Distribute data structures
 - Use per-cpu data whenever possible
 - With padding to cache lines!
- Minimize locking and expensive atomic ops
- Ideas from research have been adopted by mainstream
 - UofT/IBM Tornado/K42 projects showed techniques to improve scalability
 - Some applied to Linux scalability project
 - More recently, MIT Corey project and OSDI paper on improving Linux scalability further

