

# Lecture 7:

## Multiprocessor Scheduling

Linux scheduler illustrations from Jean-Pierre Lozi

([https://www.i3s.unice.fr/~jplozi/wastedcores/files/extended\\_talk.pdf](https://www.i3s.unice.fr/~jplozi/wastedcores/files/extended_talk.pdf))



University of Toronto, Department of Computer Science



# Multiprocessor Scheduling

---

- Why use a multiprocessor?
  - To support multiprogramming
    - Large numbers of independent processes
    - Simplified administration
    - E.g. CDF wolves, compute servers
  - To support parallel programming
    - “job” consists of multiple cooperating/communicating threads and/or processes
    - *Not* independent!
- First - the easy case: scheduling threads




# Basic MP Scheduling

---

- Given a set of runnable threads, and a set of CPUs, assign threads to CPUs
- Same considerations as uniprocessor scheduling
  - Fairness, efficiency, throughput, response time...
- But also new considerations
  - Ready queue implementation
  - Load balancing
  - Processor affinity



# Straightforward Implementation




- Scheduling events occur per CPU
  - Local timer interrupt
  - Currently-executing thread blocks or yields
  - Event is handled that unblocks thread
- Scheduler code executing on any CPU simply accesses shared queue
- What might be sub-optimal about this?




# Alternative Ready Queue Design

## Distributed Queues



## Centralized Queue



- Advantages of Distributed Queues?
- Disadvantages?



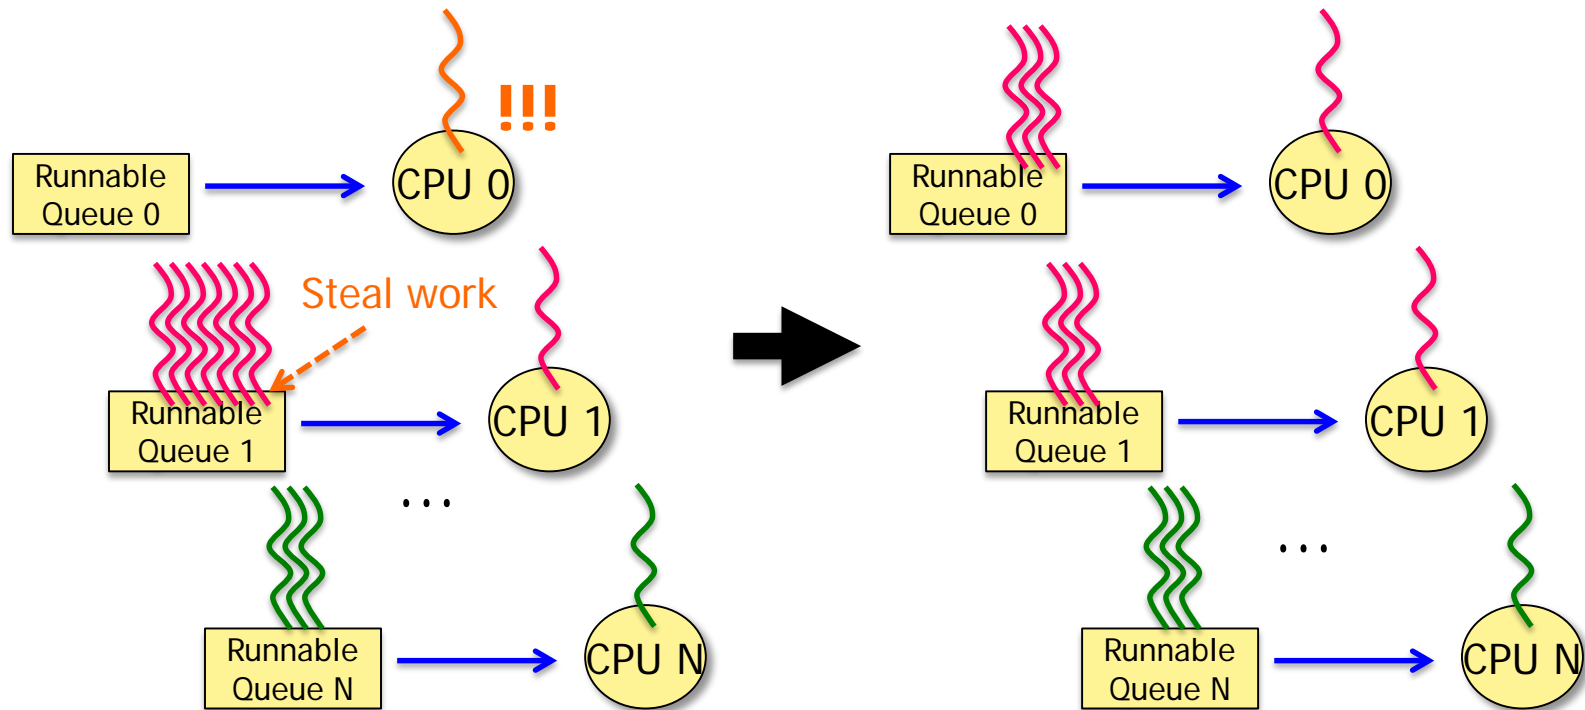
# Load Balancing

---

- Try to keep run queue sizes balanced across system
  - Main goal – CPU should not idle while other CPUs have waiting threads in their queues
  - Secondary – scheduling overhead may scale with size of run queue
    - Keep this overhead roughly the same for all CPUs
- *Push* model – kernel daemon checks queue lengths periodically, moves threads to balance
- *Pull* model – CPU notices its queue is empty (or shorter than a threshold) and steals threads from other queues
- Many systems use both



# Work Stealing with Distributed Queues



Notice a problem though?



# Processor Affinity

---

- As threads run, state accumulates in CPU cache
- Repeated scheduling on same CPU can often reuse this state
- Scheduling on different CPU requires reloading new cache
  - And possibly invalidating old cache
- Try to keep thread on same CPU it used last
  - Automatic
  - Advisory hints from user
  - Mandatory user-selected CPU
- Called “affinity scheduling”
- Do they always find a warm cache though? What can happen?





# Symbiotic Scheduling

---

- Threads load data into cache
- Expect multiple threads to thrash each others' state as they run
- Can try to detect cache needs and **schedule threads that can share nicely on same CPU**
  - Examples? What kind of threads should be scheduled together?



# Linux Scheduler: Case Study

---

*And you have to realize that there are not very many things that have aged as well as the scheduler.*

*Which is just another proof that scheduling is easy.*

*- Linus Torvalds, 2001*

Linux scheduler illustrations from Jean-Pierre Lozi

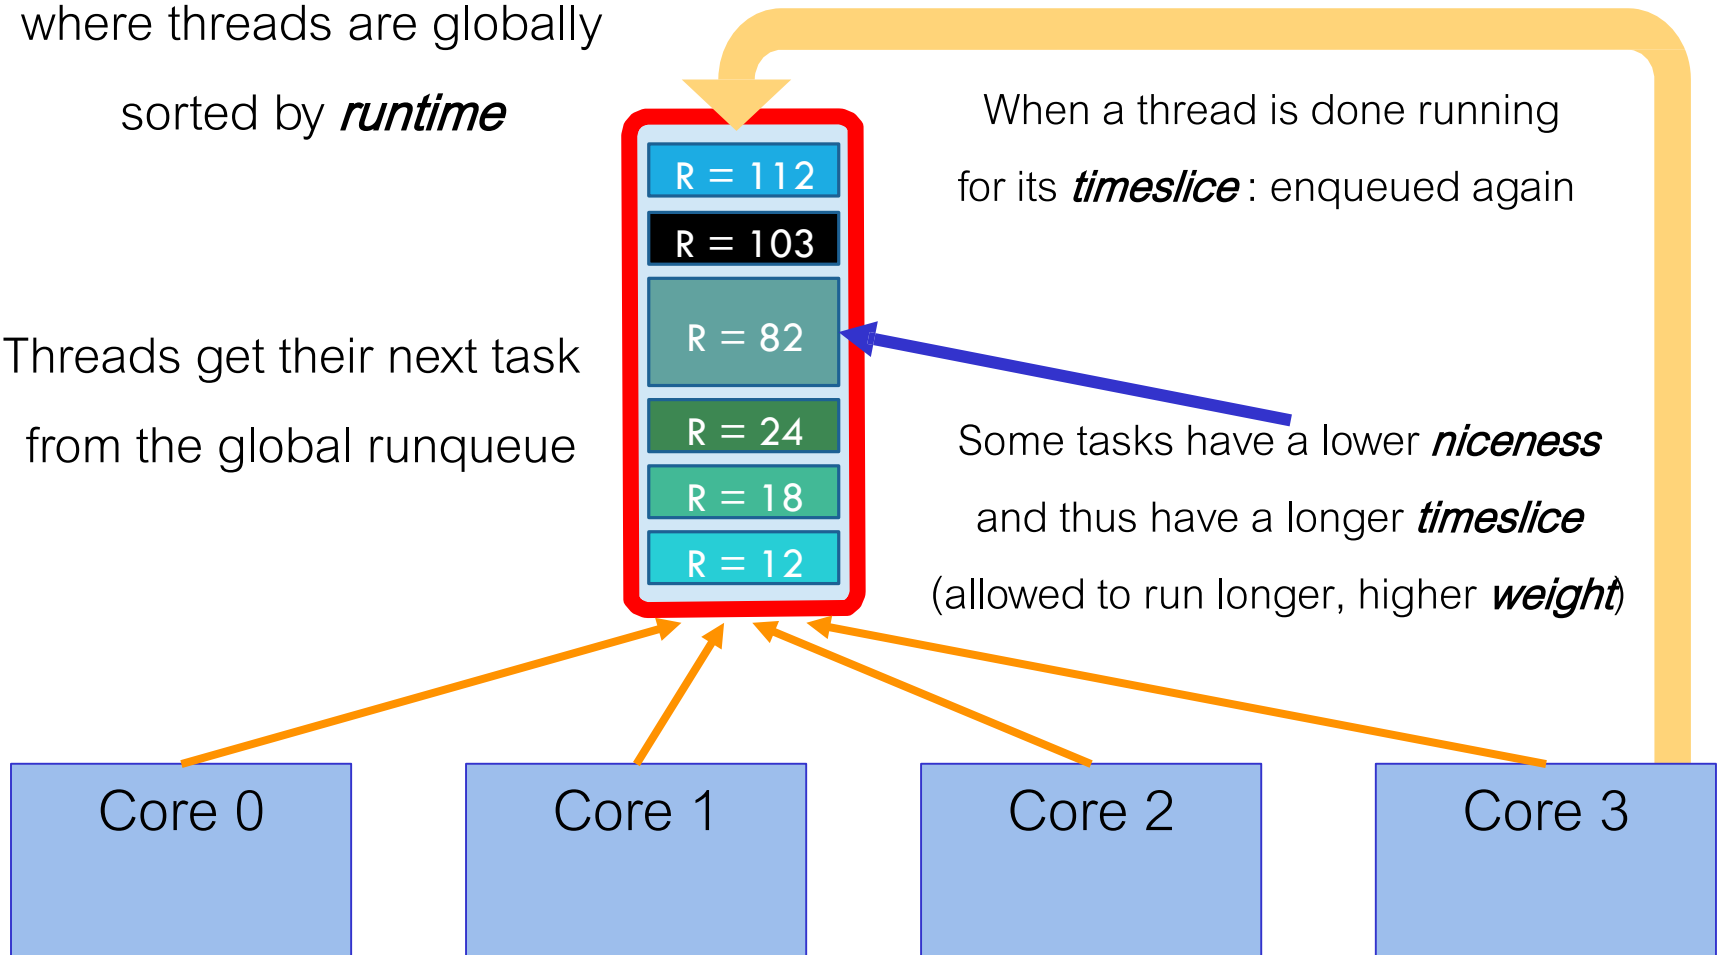
([https://www.i3s.unice.fr/~jplozi/wastedcores/files/extended\\_talk.pdf](https://www.i3s.unice.fr/~jplozi/wastedcores/files/extended_talk.pdf))



# The Completely Fair Scheduler (CFS)

Conceptually, one runqueue  
where threads are globally  
sorted by *runtime*

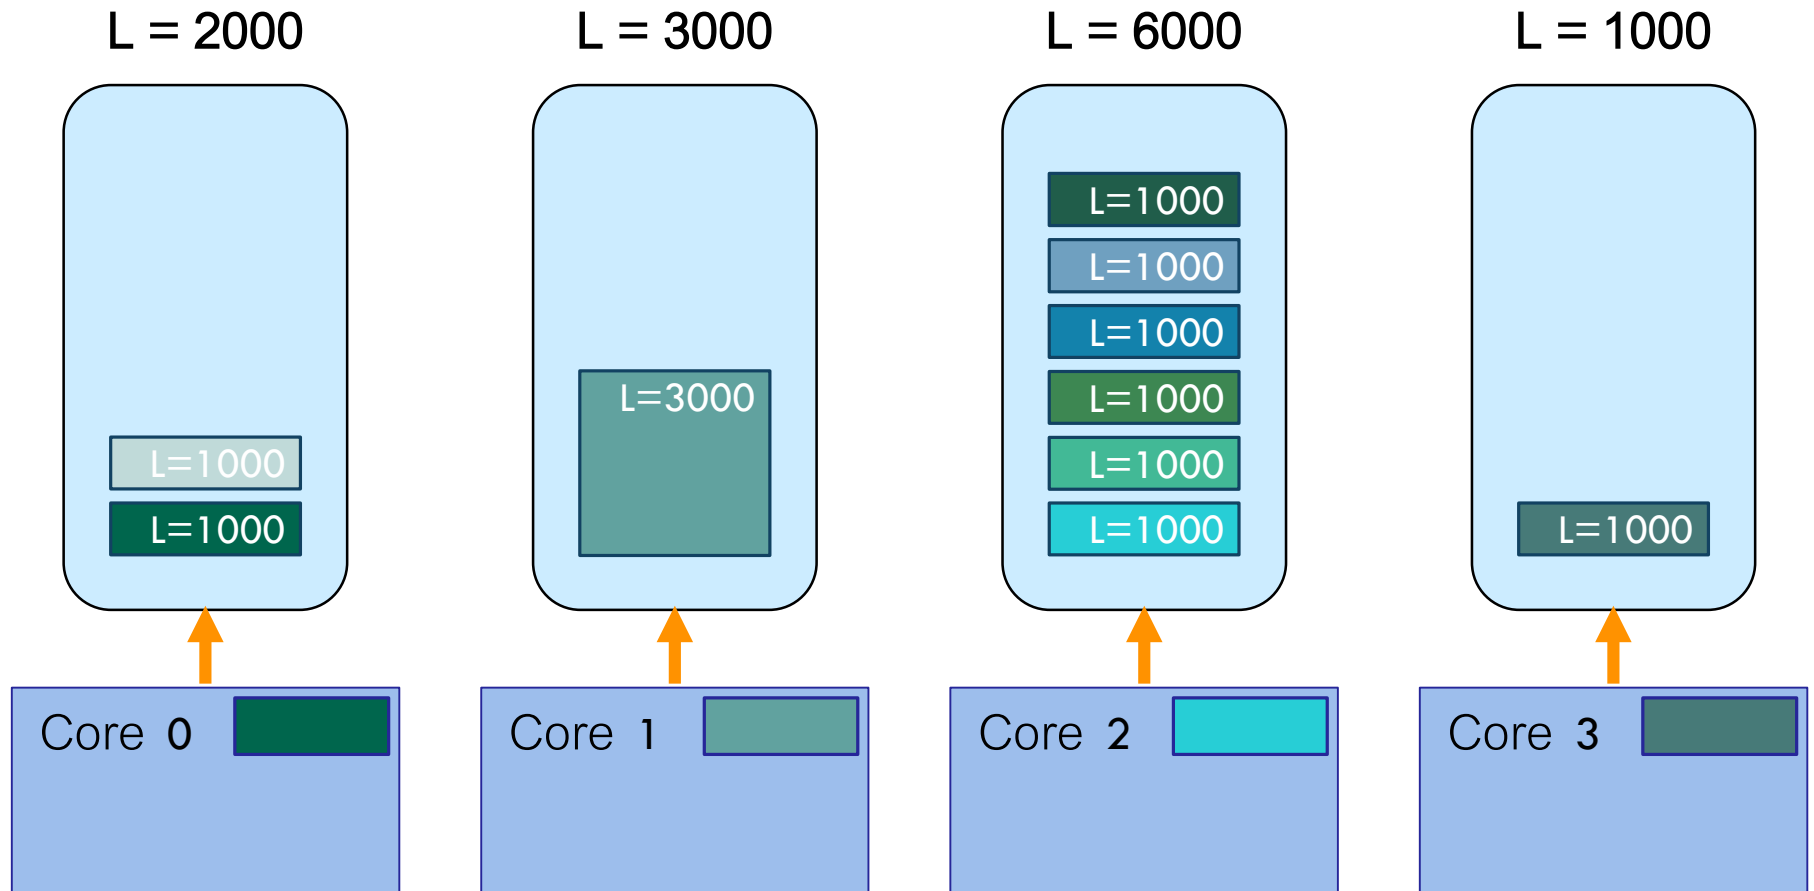
Threads get their next task  
from the global runqueue





# CFS on Multiprocessor

- Accumulated runtime is not a useful metric for load balancing
- Define CPU load of a thread:  $\text{Load} = \text{Weight} \times \% \text{CPU}$





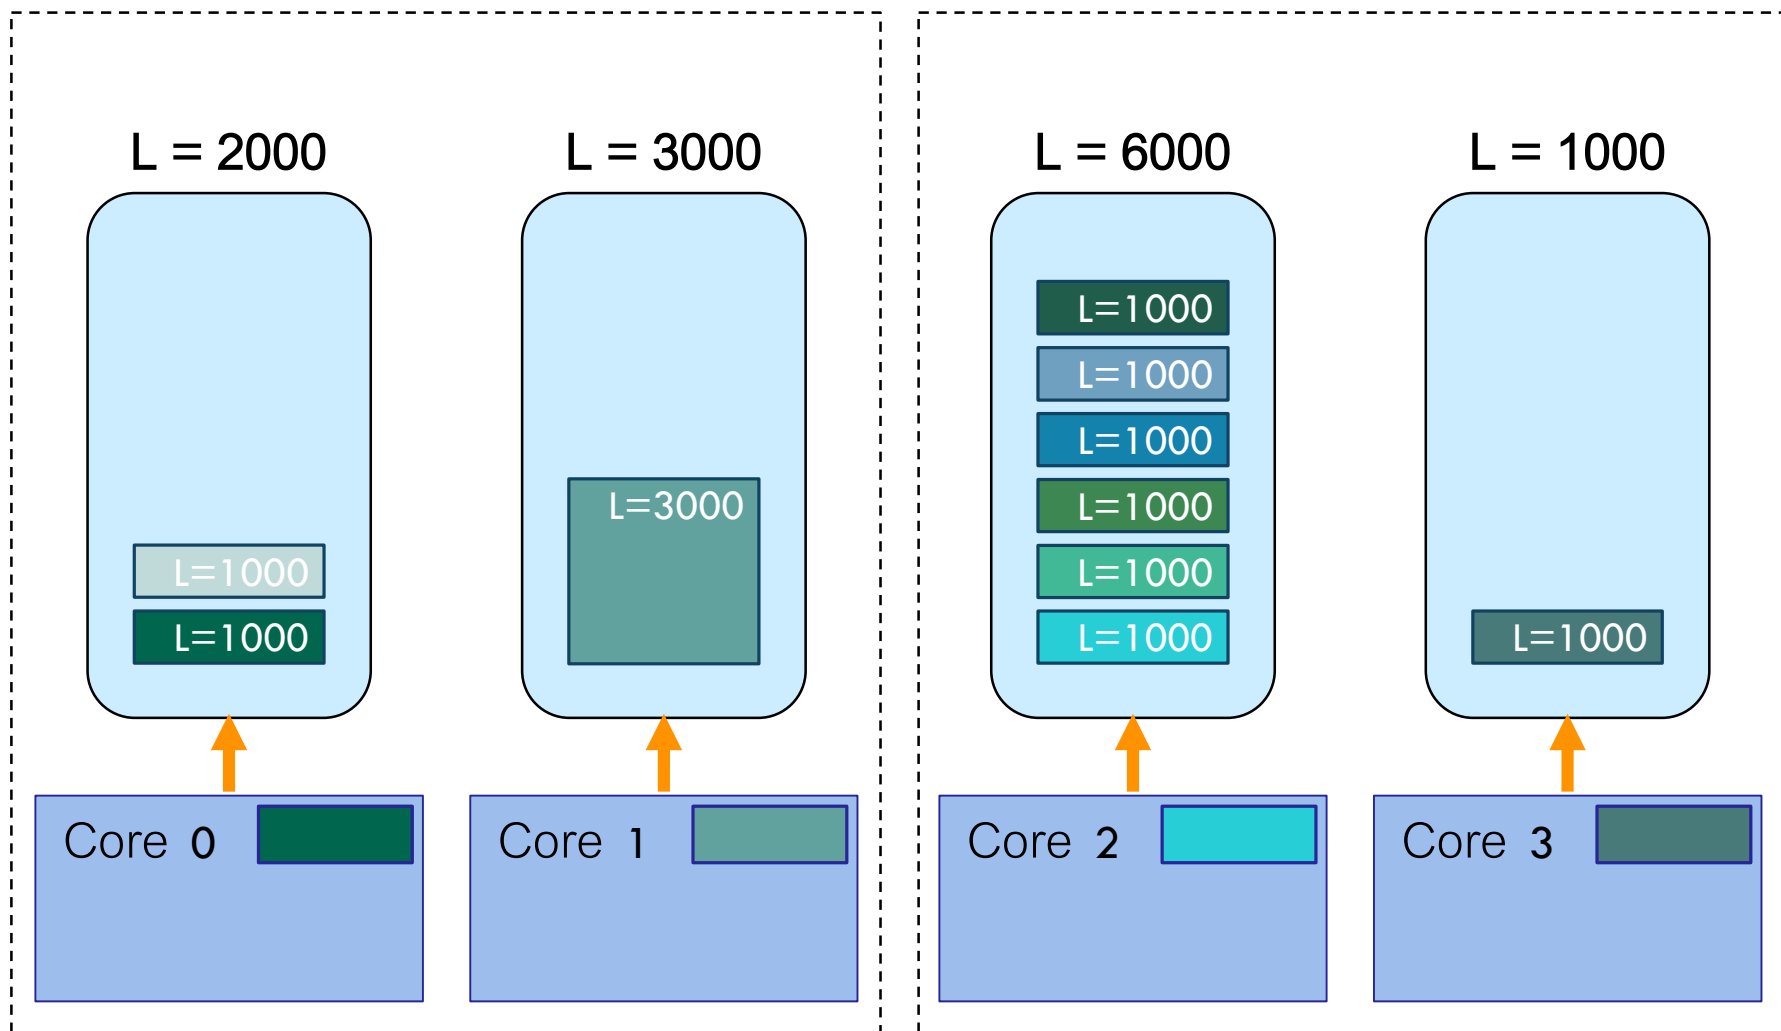
# Linux Scheduler Basics

---

- per-CPU runqueues
  - Actually a red-black tree with the CFS scheduler
- CPUs are organized into *scheduling domains*
  - A set of CPUs whose load is kept balanced by kernel
  - Each domain contains a set of groups
- Load balancing is hierarchical
  - On each tick, recomputes local load statistics
  - Checks if time to invoke `load_balance()` for each domain from base to top-level

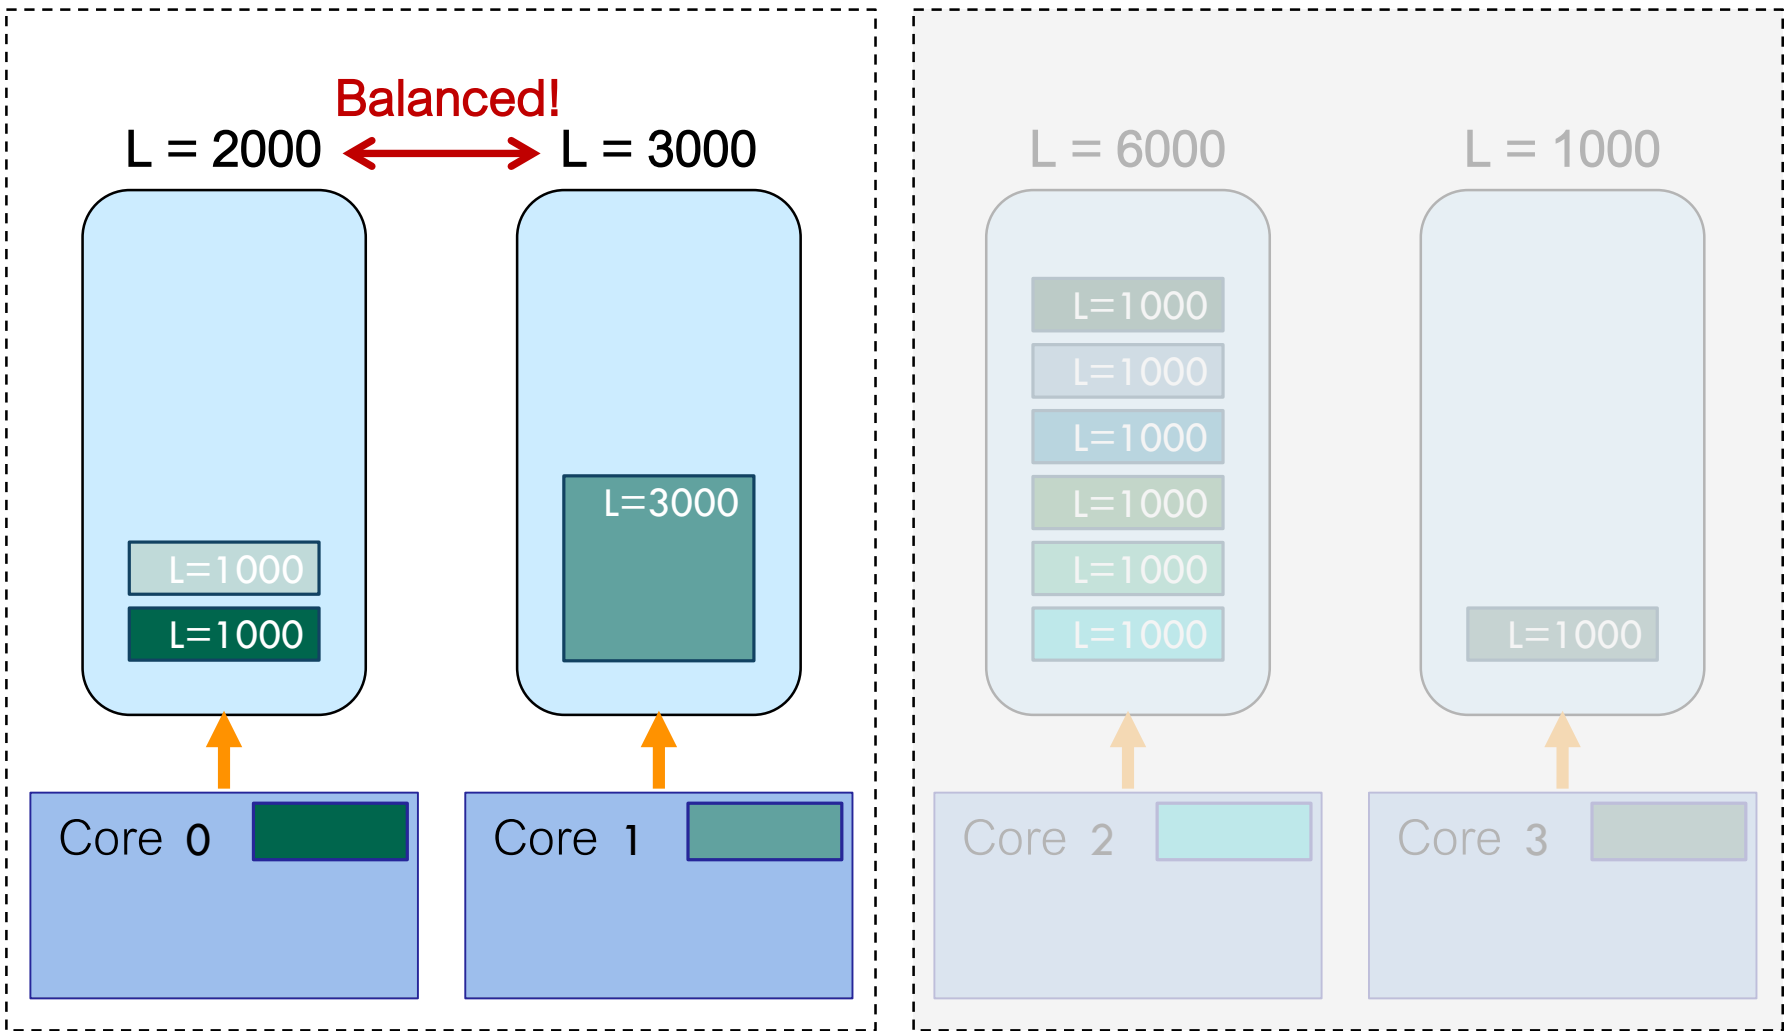


# Hierarchical Load Balancing



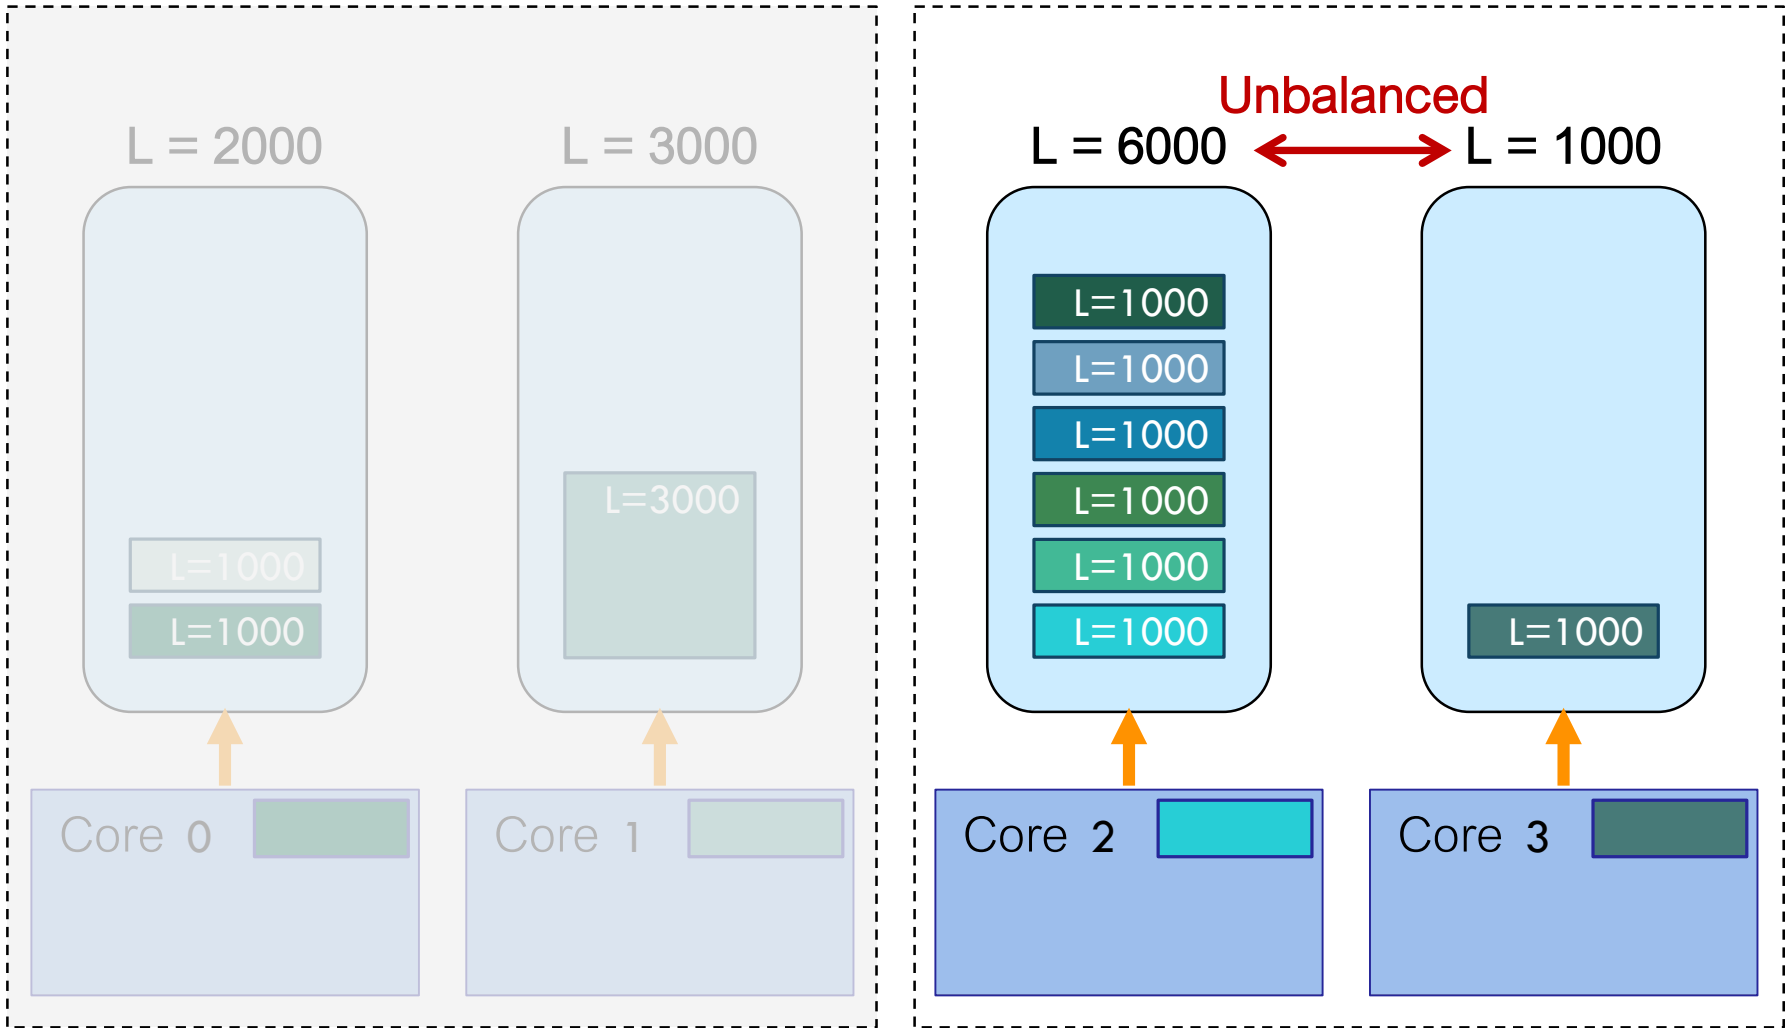


# Hierarchical Load Balancing





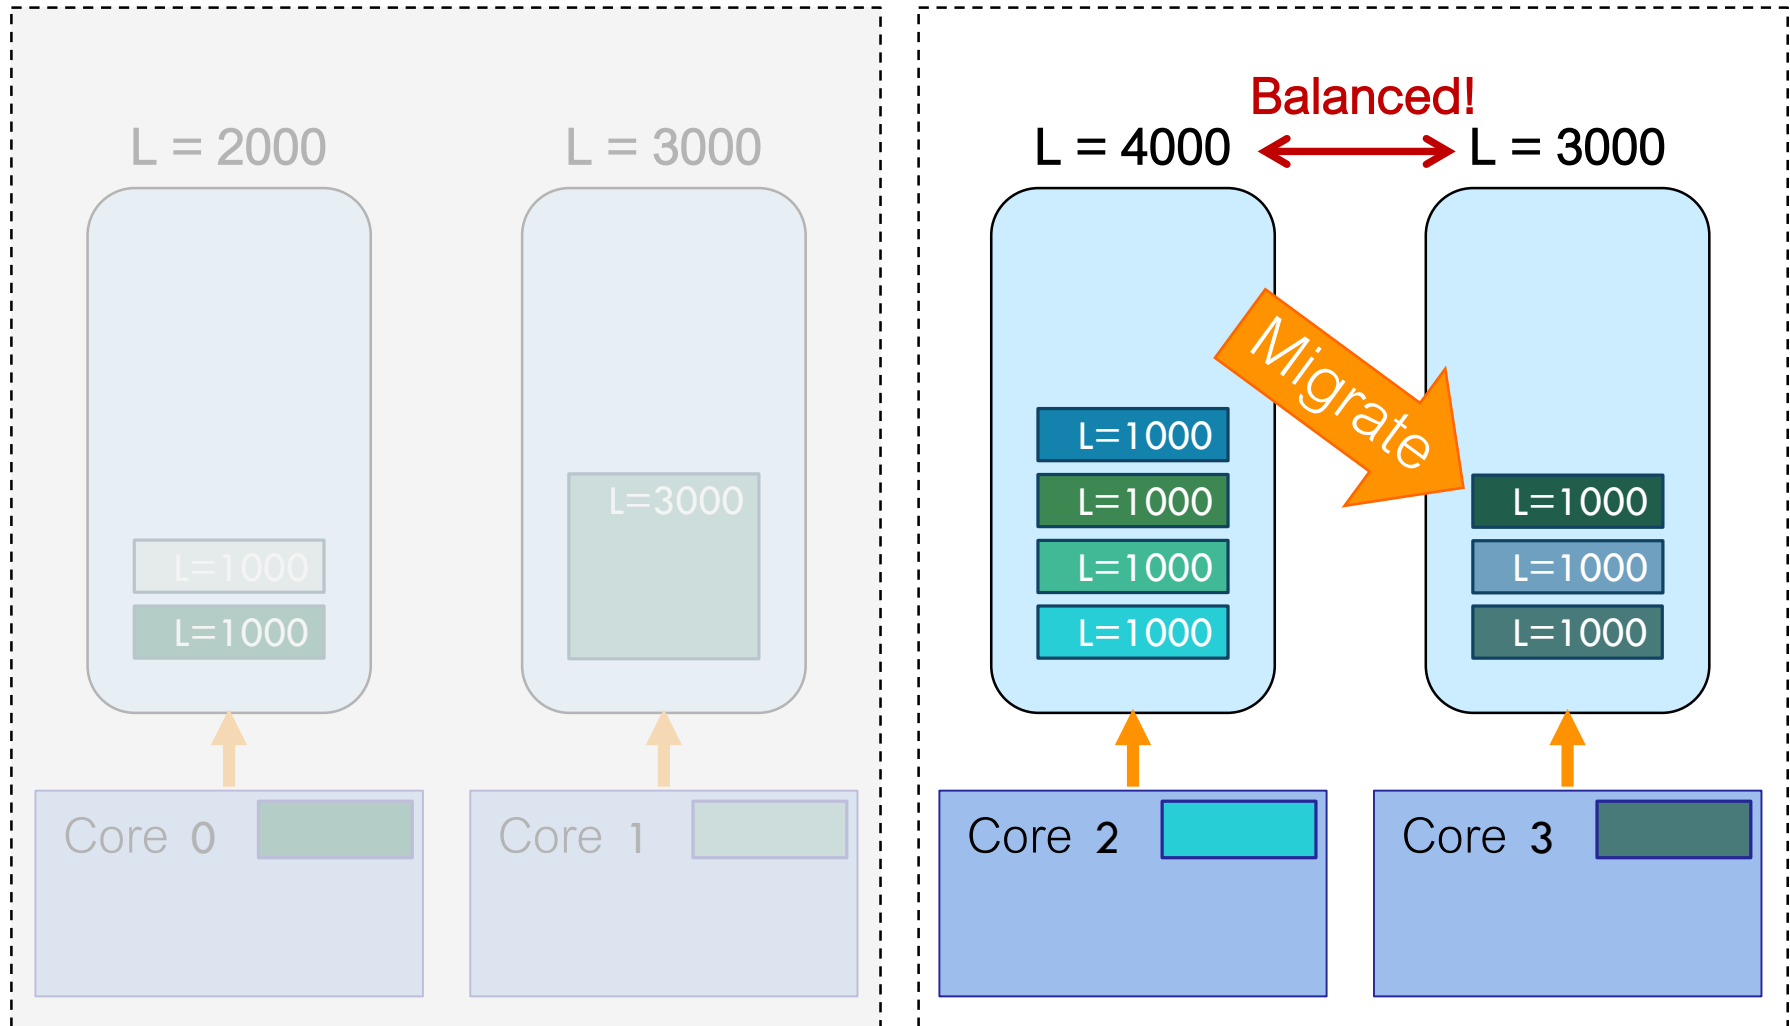
# Hierarchical Load Balancing







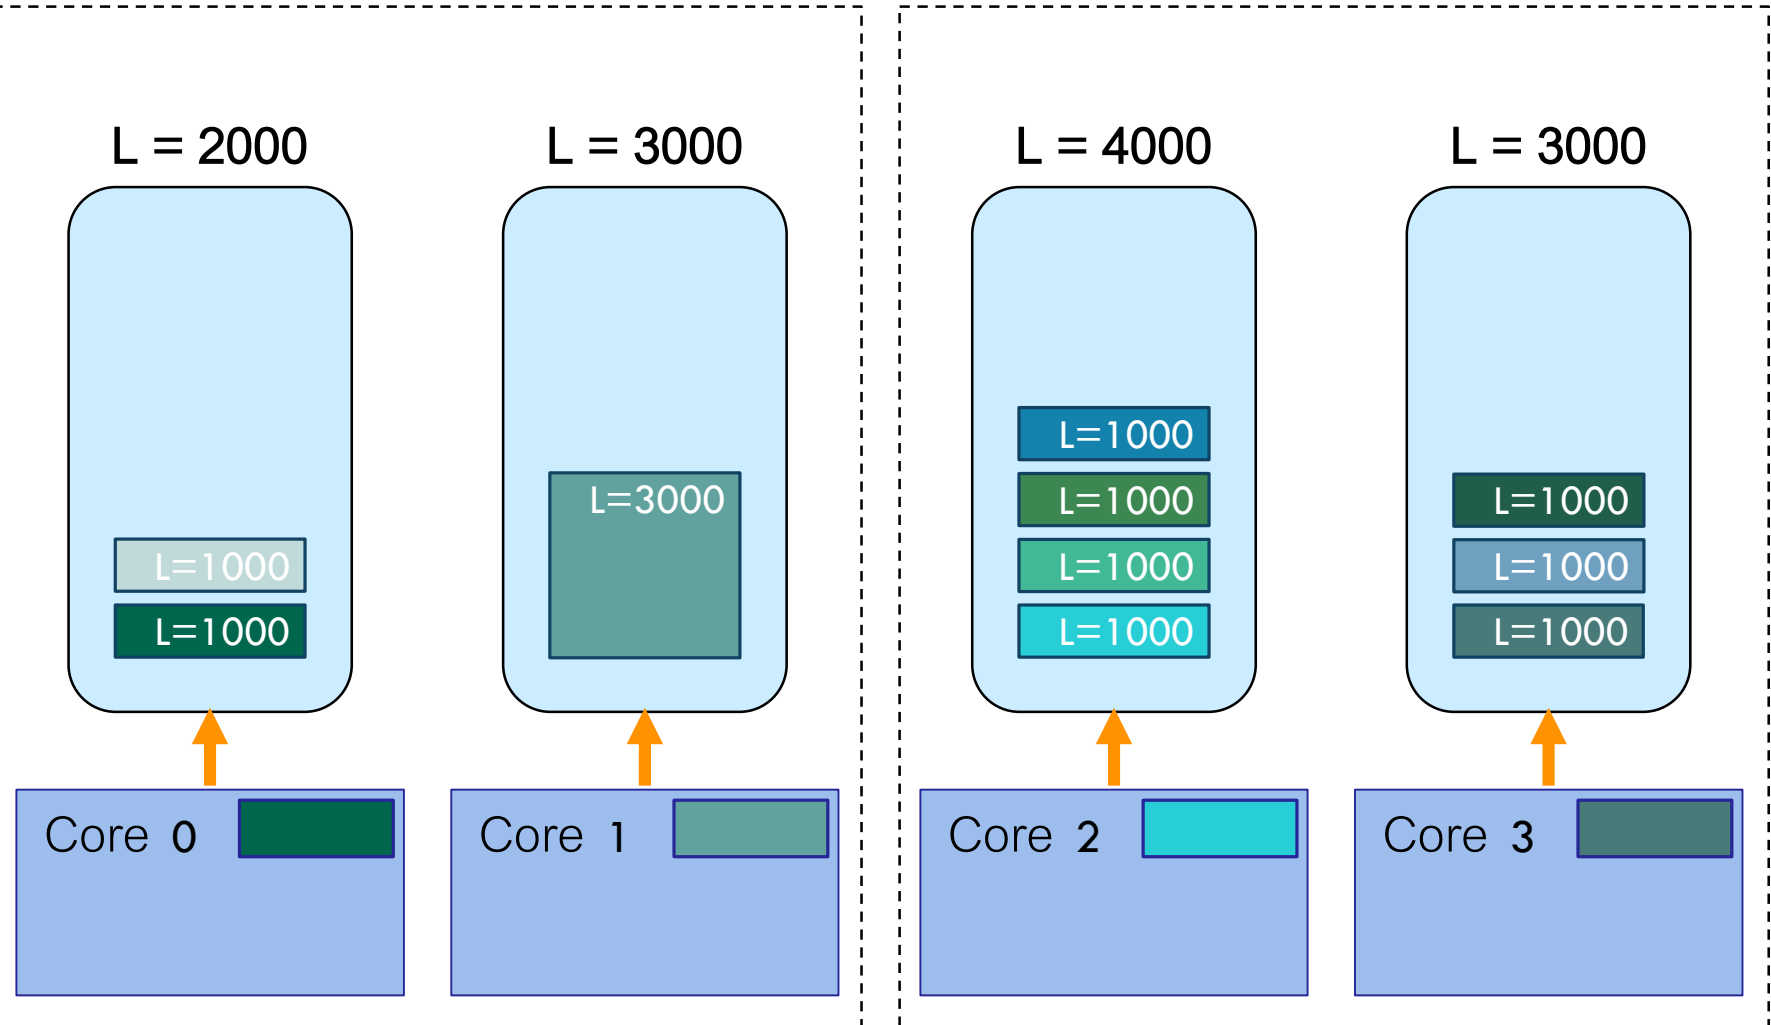
# Hierarchical Load Balancing





# Hierarchical Load Balancing

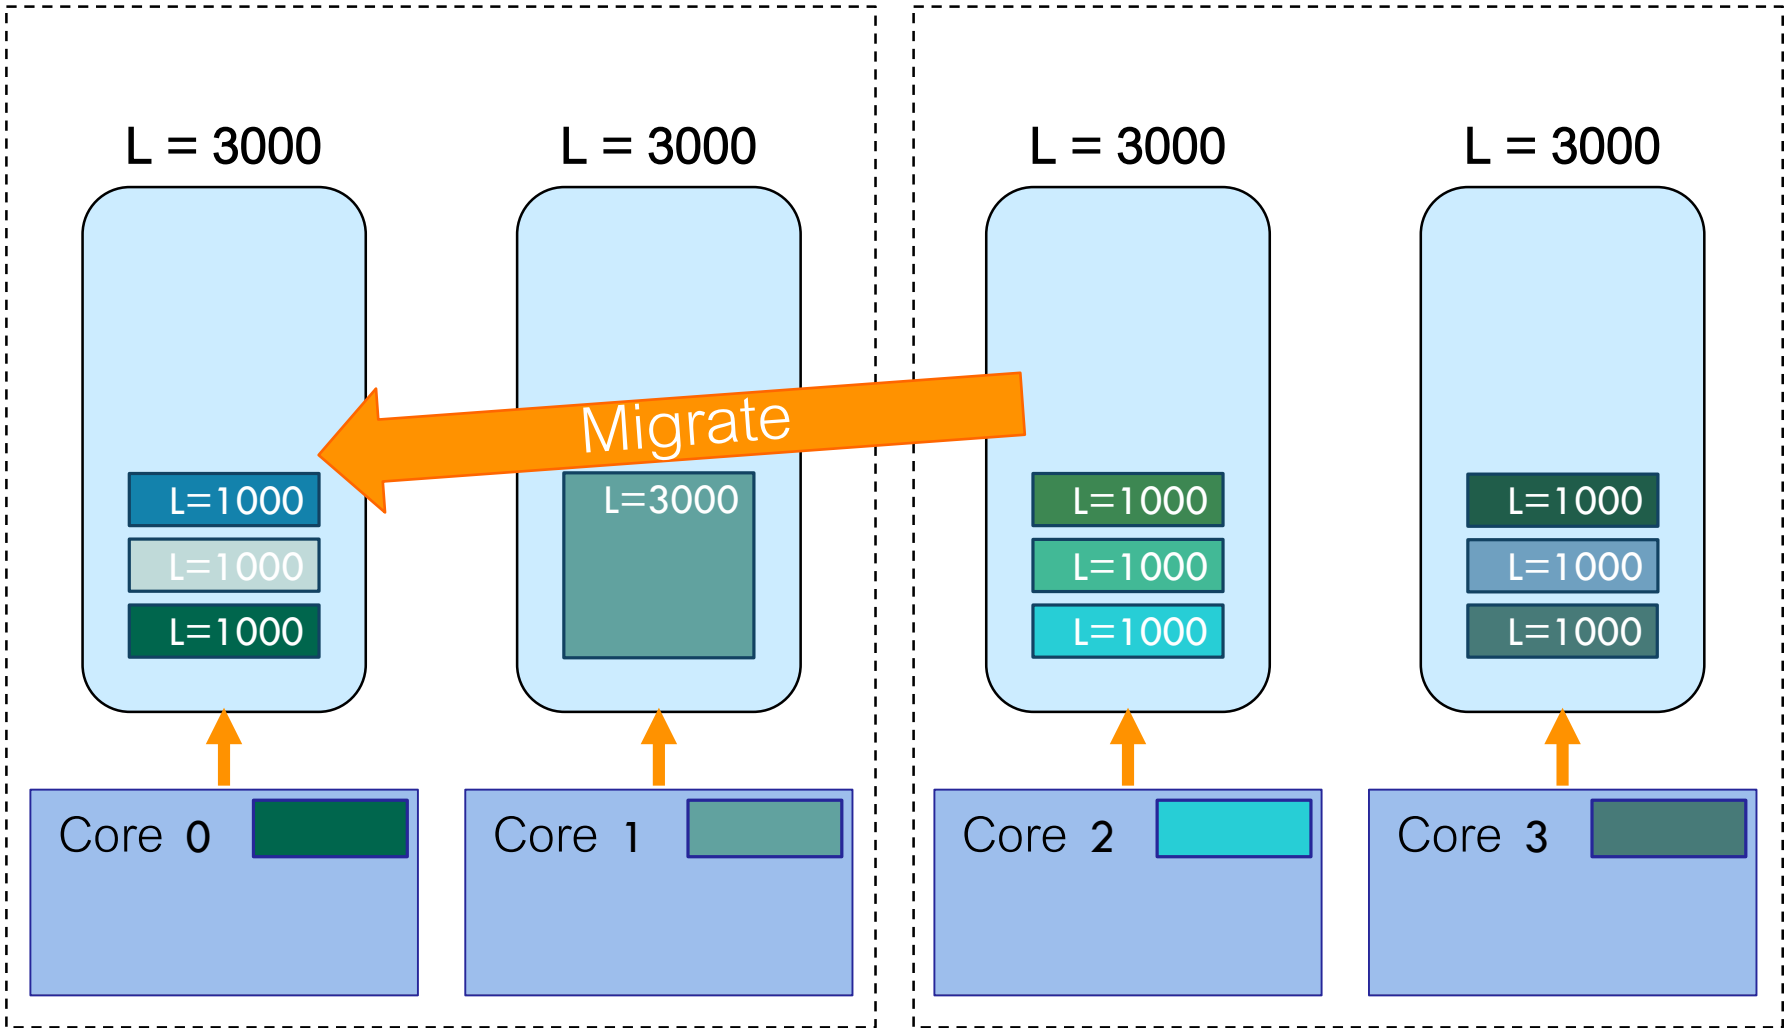
Average(L) = 2500 ↔ Unbalanced ↔ Average(L) = 3500





# Hierarchical Load Balancing

Average(L) = 300  $\xleftrightarrow{\text{Balanced}}$  Average(L) = 3000





# Linux Load Balancing

---

- Finds busiest group in current domain
- Finds busiest queue (CPU) in that group
- Invokes `move_tasks` to actually move threads
  - `move_tasks` attempts to preserve affinity when finding task to move
    - Can't be currently executing
    - Target CPU must be allowable for task
    - Target CPU is idle OR process is not "cache hot" OR kernel has failed repeatedly to move processes
- Actual load calculations in Linux are quite complex
  - Can fail to achieve balance and fairness in some scenarios
- Further reading (optional):
  - The Linux Scheduler – a decade of wasted cores (Eurosyst 2016)
  - <https://www.ece.ubc.ca/~sasha/papers/eurosyst16-final29.pdf>




# Parallel Job Scheduling

---

- “Job” is a collection of processes/threads that cooperate to solve some problem (or provide some service)
  - *Not* independent!
- How the components of the job are scheduled has a major effect on **performance**
  - Want scheduler to be **aware of dependences**
- We will look at two major strategies
  - **Space sharing** – each job has dedicated processors
  - **Time sharing** – multiple jobs share same processors



# Why Job Scheduling Matters?

- Recall threads in a job are not independent 
  - Synchronize over shared data
    - De-schedule lock holder, other threads in job may not get far
  - Cause/effect relationships (e.g. producer-consumer problem)
    - Consumer is waiting for data on queue, but producer is not running
  - Synchronizing phases of execution (barriers)
    - Entire job proceeds at pace of slowest thread



# Forms of scheduler-awareness

---


- 1. Know threads are related, schedule all at same time
  - Space sharing: all threads are from same job
  - Time sharing: group threads that should be scheduled together
- 2. Know when threads hold spinlocks
  - Don't deschedule lock holder
  - Extends timeslice, but not indefinitely
- 3. Know about general dependences
  - E.g. infer producer/consumer relationships




# Space Sharing Scheduling

- Divide processors into groups
  - Fixed, variable, or adaptive
- Assign job to dedicated set of processors
  - Ideally one CPU per thread in job
- Job waits until required number of CPUs are available (batch scheduling)


**Fixed:**  
Always 2  
groups of  
4 CPUs.



**Variable:**  
Currently 3  
groups of 2,  
4, and 2  
CPUs;  
changes as  
jobs come  
and go.



**Adaptive:**  
Job can ask  
for more  
CPUs as it  
runs.







# Space Sharing

---

- Typically used on supercomputers
- Pros:
  - All runnable threads execute at the same time
  - Reduce context switch overhead (no involuntary preemption)
  - Strong affinity
- Cons?
  - Inflexible
    - CPUs in one partition may be idle while another partition has multiple jobs waiting to run
  - Difficult to deal with dynamically-changing job sizes
    - Adaptive scheme is complicated and uncommon



# Choosing Jobs to Run

---

- At job creation, specify number of threads
- Scheduler finds set of CPUs
  - May negotiate with application
- How should scheduler choose jobs to assign to CPUs?

What is optimal (in terms of average wait time)?

- Uniprocessor scheduling → Shortest Job First (SJF) (shortest expected next CPU burst)
- MP version → smallest expected number of CPU cycles  
(cycles == num\_cpus \* runtime)



# Estimating Runtime


---

- Estimates typically come from users who submit the jobs
  - Low estimates make it “easier” to do scheduling
  - But cause trouble if not accurate!
  - Soln: kill jobs that exceed estimate
    - What user behaviour does that incentivize?
- How accurate are user estimates?
- Can automatic estimates based on history do better?
- How much does it matter?



# Space Sharing - FCFS

- Scheduling convoy effect
  - Long average wait times due to large job
  - Exists with FCFS uniprocessor batch systems
  - Much worse in parallel systems
    - Fragmentation of CPU space






# Solution: Backfilling

- Fill CPU “holes” from queue in FCFS order
- Not FCFS anymore. What can happen?
- Want to prevent “fill” from delaying threads that were in queue earlier
  - EASY (E<sub>xtensible</sub> A<sub>rgonne</sub> S<sub>cheduling</sub> s<sub>Y</sub>stem)
  - Make reservation for next job in queue

Scheduler queue (CPUs, time)

E	D	C	B	A
1,3	2,2	4,2	2,2	1,4





# Variations on Backfilling

---

- EASY
  - 1. Used FCFS to order jobs in queue
  - 2. Made reservation for first blocked job in queue
  - 3. Backfilled jobs by looking at queue one at a time
- 1. **Ordering alternative**: include priority in queue
  - Administrative to distinguish between users
  - User to distinguish between own jobs
  - Scheduler to prevent starvation
- 2. **Reservation alternatives**
  - All queued jobs get a reservation (too much can go wrong)
  - Queued job gets a reservation if it has been waiting more than a threshold
- 3. **Queue lookahead**
  - Use dynamic programming to determine optimal packing



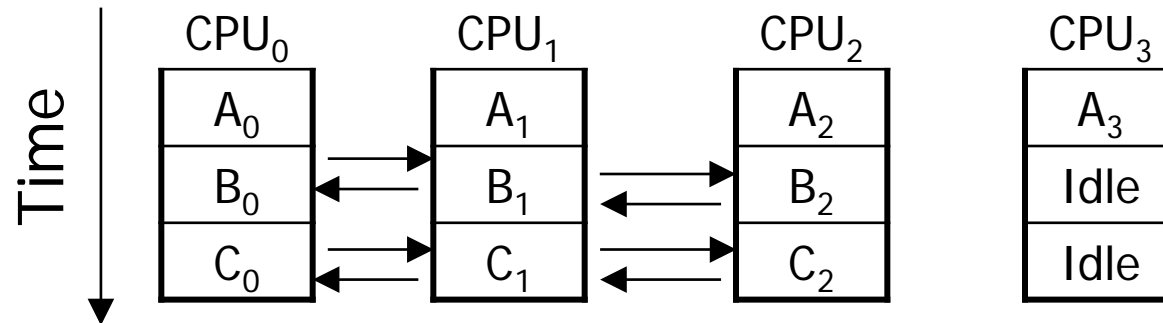
# Parallel Time Sharing

---

- Each CPU may run threads from multiple jobs
  - But with awareness of jobs
- Co-scheduling (Ousterhout, 1982)
  - Identify “working set” of processes (analogous to working set of memory pages) that need to run together
- Gang scheduling
  - All-or-nothing → co-scheduled working set is all threads in the job
  - Get scheduling benefits of dedicated machine
  - Allows all jobs to get service



# Gang Scheduling Example



- Multiprogramming level is typically controlled by either:
  - Monitoring memory demand, or
  - Fixed number of slots (rows)
    - e.g. IBM LoadLeveler Gang Scheduling allows up to 8 sets of jobs to be multiprogrammed on a set of CPUs





# Gang Scheduling Issues


---

- All CPUs must context switch together
  - To avoid fragmentation, construct groups of jobs that fill a slot on each CPU
    - E.g., 8-CPU system, group one 4-thread job with two 2-thread jobs
  - Inflexible
    - If 4-thread job blocks, should we block entire group, or schedule group and leave 4 CPUs idle?
- Alternative 1: Paired gang scheduling
  - Identify groupings with complementary characteristics and pair them.  
When one blocks, the other runs
- Alternative 2: Only use gang scheduling for thread groups that benefit
  - Fill holes in schedule with any single runnable thread from those remaining



# Example: Effect of Gang Scheduling

- LLNL gang scheduler on 12-CPU Digital Alpha 8400
  - Parallel gaussian elimination program
  - Run concurrently with 12 single-threaded interfering processes
  - Benefits due to synchronization effects and better cache use



Source:

“Expanding symmetric multiprocessor capability through gang scheduling”,  
Morris A. Jette, in Job Scheduling Strategies for Parallel Processing, LNCS 1998, Volume 1459/1998



# Forms of scheduler-awareness

---

- ✓ 1. Know threads are related, schedule all at same time
  - Space sharing: all threads are from same job
  - Time sharing: group threads that should be scheduled together
  
- 2. Know when threads hold spinlocks
  - Don't deschedule lock holder
  - Extends timeslice, but not indefinitely
  
- 3. Know about general dependences
  - E.g. infer producer/consumer relationships



# Knowing about Spinlocks

---

- Thread acquiring spinlock sets **kernel-visible flag**
- Clears flag on release
- Scheduler will **not immediately deschedule** a thread with the flag set
  - Gives thread a chance to complete critical section and release lock
  - Spinlock-protected critical sections are (supposed to be) short
  - Does not defer scheduling indefinitely



# Forms of scheduler-awareness

---

- ✓ 1. Know threads are related, schedule all at same time
  - Space sharing: all threads are from same job
  - Time sharing: group threads that should be scheduled together
  
- ✓ 2. Know when threads hold spinlocks
  - Don't deschedule lock holder
  - Extends timeslice, but not indefinitely
  
- 3. Know about general dependences
  - E.g. infer producer/consumer relationships



# Knowing General Dependences

---

- Implicit Co-scheduling (Arpaci-Dusseau et al.)
- Designed for workstation cluster environment
  - Explicit messages for all communication/synchronization
  - MUCH more expensive if remote process is not running when local process needs to synchronize
- Communicating processes decide when it is beneficial to run
  - Infer remote state by observing local events
    - Message round-trip time
    - Message arrival
- Local scheduler uses communication info in calculating priority



# OS Noise

---

- Or: how to schedule OS activities
- Massively parallel systems are typically split into I/O nodes, management nodes, and compute nodes
  - Compute nodes are where the real work gets done
  - Run customized, lightweight kernel on compute nodes
  - Run full-blown OS on I/O nodes and mgmt nodes
  - Why?
- Asynchronous OS activities perturb nice scheduling properties of running jobs together
  - Up to a factor of 2 performance loss in real large-scale jobs
  - Need to either eliminate OS interference, or find ways to coordinate it as well



# Illustration of OS Noise Issue

---


- The following 4 slides are extracted from a tutorial given at EuroPar 2004 ([www.di.unipi.it/euopar04/Tutorial2/tutorial\\_euopar04.ppt](http://www.di.unipi.it/euopar04/Tutorial2/tutorial_euopar04.ppt))
    - They illustrate the issue with OS Noise on the ASCI Q supercomputer, published in Supercomputing 2003 (“The case of the missing supercomputer performance”)
- “Achieving Usability and Efficiency in Large-Scale Parallel Computing”
- Kei Davis and Fabrizio Petrini {kei,fabrizio}@lanl.gov
  - Performance and Architectures Lab (PAL), CCS-3
    - (Computer and Computational Sciences Division, Los Alamos National Labs)





# Performance of SAGE on 1024 nodes

- Performance consistent across QA and QB (the two segments of ASCI Q, with 1024 nodes/4096 processors each) ← *4 processors / node*
  - Measured time 2x greater than model (4096 Processor Elements)



There is a difference  
why ?


*OS activity is the culprit!*

Lower is better!




# The effect of the noise

- An application is usually a sequence of a computation followed by a synchronization (collective):




- But if an event happens on a single node, it can affect all the other nodes






# Effect of System Size



- The probability of a random event occurring increases with the node count



# Tolerating Noise: Buffered Coscheduling



- We can tolerate the noise by coscheduling the activities of the system software on each node



# Example Cluster Scheduler: SLURM

---

- Simple Linux Utility for Resource Management
  - Performs resource management within single cluster => 3 roles:
    - **Allocates** access to computer nodes and their interconnect
    - **Launches** parallel jobs and **manages** them (I/O, signals, time limits, etc.)
    - **Manages contention** in the queue
- Developed by Lawrence Livermore National Lab (LLNL)
  - With help from HP, Bull (European high performance computing company), Linux NetworX, and others
  - Open Source
  - Extensible, provides flexible plugin mechanism
  - Active development still on-going
- Widely used on high performance compute clusters



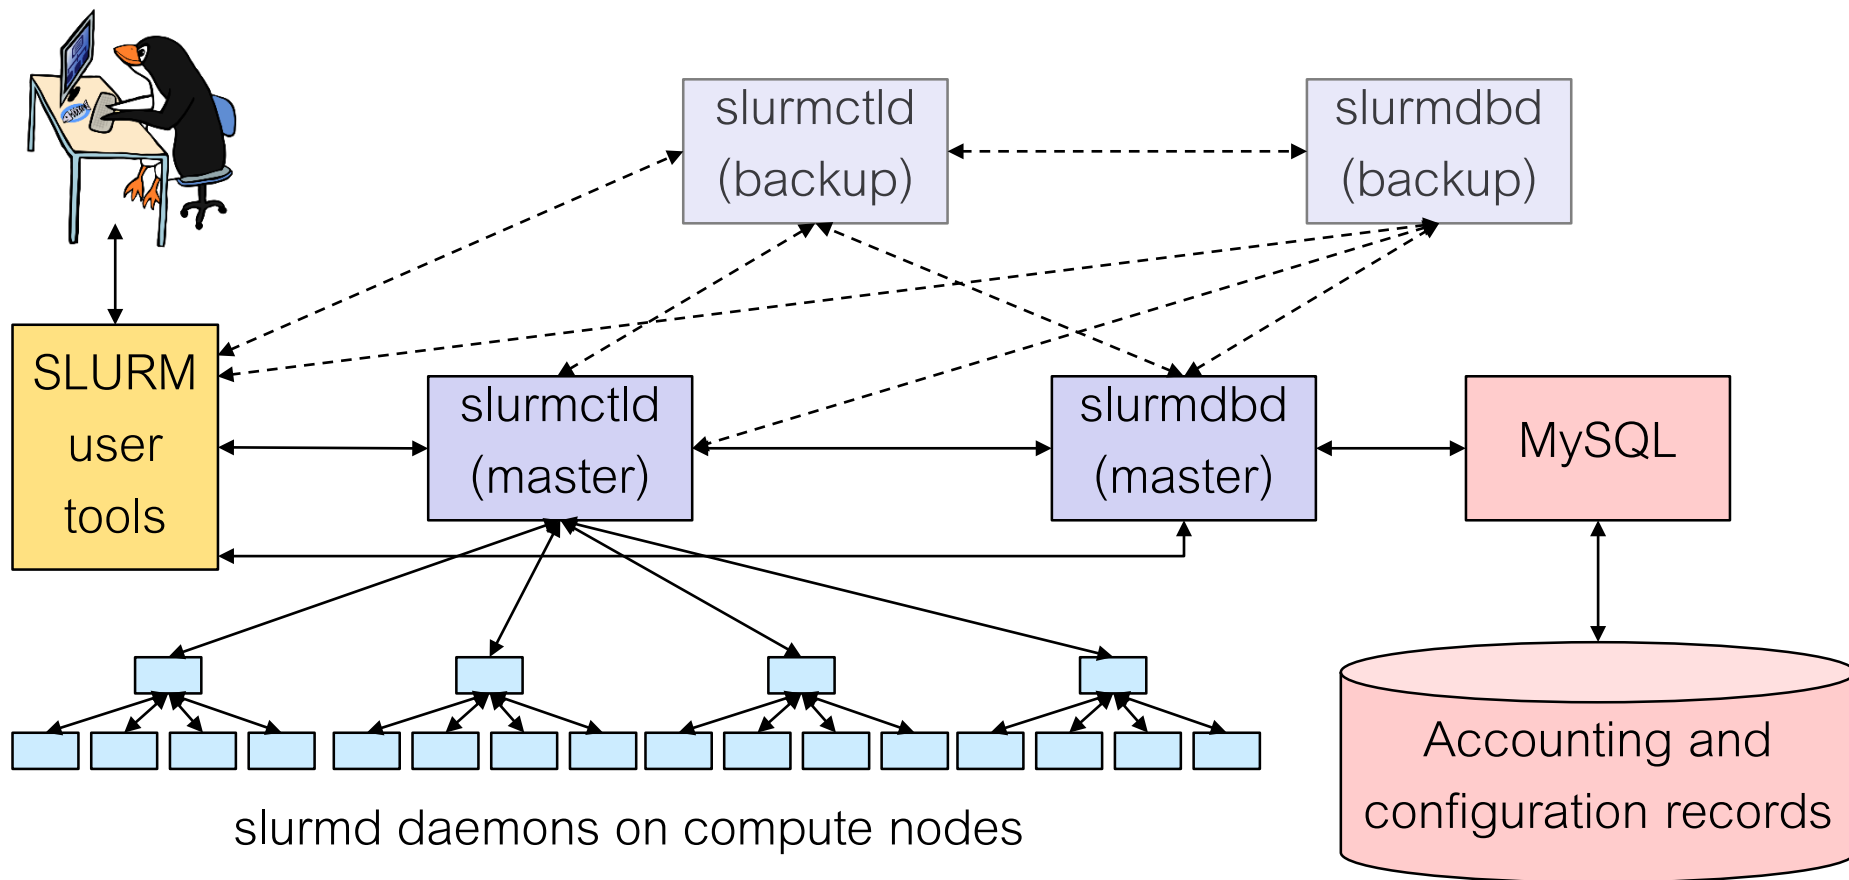
# SLURM Features

---

- Plugins support multiple scheduling policies
  - FIFO
  - Backfilling
  - Gang Scheduling
    - Requires multi-core awareness at slurmctld
  - Priority-based preemption
  - Topology-aware scheduling
    - Reduce contention on interconnect
- Includes many management & accounting features



# SLURM Architecture



slurmd daemons on compute nodes

(hierarchical communications with configurable fanout)

Copyright 2017 SchedMD LLC

<http://www.schedmd.com>

Primary communication

Backup communication



# Further readings

---

- Scheduling problem encountered in many contexts (e.g., datacenters)
- Omega (Eurosys 2013)
  - <https://research.google.com/pubs/pub41684.html>
  - [https://people.csail.mit.edu/malte/pub/talks/2013-04-17\\_eurosys-omega.pdf](https://people.csail.mit.edu/malte/pub/talks/2013-04-17_eurosys-omega.pdf)
- Hawk (USENIX 2015)
  - <https://www.usenix.org/conference/atc15/technical-session/presentation/delgado>
- Mesos (NSDI 2011)
  - <https://people.eecs.berkeley.edu/~alig/papers/mesos.pdf>
- Sparrow (SOSP 2013)
  - <https://dl.acm.org/citation.cfm?id=2522716>
- Thoth (VLDB 2013)
  - <https://dl.acm.org/citation.cfm?id=2733062>
- And many more ...





# Announcements

---

- Midterm
  - Friday!
  - Bring T-card!
  - Covers up end of last week (OS Scalability, not Scheduling)
- Location and logistics:
  - Starts at **9AM**, as discussed in the first week of classes
  - Please be on time, we start at **9:10 sharp** => 110 minutes
  - Check website and Piazza announcements!