

Week 12b: Security

CSC 469 / CSC 2208

Fall 2018



University of Toronto, Department of Computer Science



What it's all about

- Managing risks. Loss of ...
 - Confidentiality/Privacy, Integrity, Availability/Access
 - Risk Analysis: $\text{cost/loss} * \text{loss freq.}$ vs. cost to protect
 - Engineering trade-offs, not either-or decisions
 - often, $\text{Security} == 1 / (\text{functionality} * \text{convenience})$
- Vulnerabilities
 - examples abound, many reasons behind
- Countermeasures
 - carefulness, cryptography, firewalls, detection, recovery



Some key security goals

- Confidentiality: keep information content away from the unauthorized
- Integrity: prevent undetected, unauthorized modification of data
- Availability: ensure that resources and services are available when needed
- Authentication: prove the identity of entities or source of information
- Non-repudiation: prevent denial of previous commitments
- Privacy and Anonymity: protect personally identifiable info



Some key security problems

- 1. Misplaced trust
- 2. Buggy implementations
- 3. Poor configuration choices
- ...
- 12. Unsafe design assumptions
- ...
- 997. Cryptanalysis

(<https://xkcd.com/538/>)



Terminology: Threats

- Threat: A potential vector (means, mechanism) for a system's security to be compromised
 - An attack exercises a threat
 - A successful attacks leads to a security compromise
- Examples of threats:
 - Network traffic arriving from the internet
 - Self-administered systems connected to a corporate (or university) LAN



Terminology: Vulnerabilities

- A vulnerability is a flaw in a system that has a security implication
- Examples:
 - Unchecked string copy allows buffer overflow
 - Administrator forgets to disable debug mode on a program during configuration, leaving unsafe but convenient features in deployed service
 - Naïve home user buys wireless router, but does not alter default password on router
- Compromises occur when an attacker matches threats with vulnerabilities



Scary aspects of “bad guys”

- Patience and time
 - Historically successful crackers have been willing to spend endless hours trying to get into systems
- Automated Tools
 - Crackers don't even have to know anything anymore
 - Copious “cookbooks” and packaged kits
 - One clever person finds a hole, everyone runs her tools
- New profit motive
 - Rent-a-bot-net brokers



IRC Hacker Market

2006
data





Figure 15: Price schedule for compromised hosts.

- From “An Inquiry into the Nature and Causes of the Wealth of Internet Miscreants”, Franklin et al., in Proc. Computer & Communications Security (CCS), October 2007



IRC Hacker Market



- From “An Inquiry into the Nature and Causes of the Wealth of Internet Miscreants”, Franklin et al., in Proc. Computer & Communications Security (CCS), October 2007



...and if that weren't bad enough

- attackers only need one weakness
 - no need to break thru strongest wall
 - they'll try lots and exploit the weakest



Early Example: The Morris Worm

- Released on November 2, 1988
- Written by Robert T. Morris
- Invaded around 6,000 computers within hours (10% of the Internet at the time)
- Disabled many systems and services
- Morris had a friend post instructions for disabling the worm - but it was too late
- Damage estimates between \$10,000 and \$97 million (shows how hard it is to estimate costs)
- Details in June 1989 Comm. of the ACM “Crisis and Aftermath”, Eugene H. Spafford



How the worm worked

- Copied itself to remote systems via 3 holes
- Exploit hole in finger daemon that caused buffer overflow to create remote shell
 - `gets()` used to read input
- Exploit hole in Unix sendmail daemon
 - `listen()`'s on TCP port, `accept()`'s connections from mailers
 - Exchanges messages about mail envelope and content
 - when running in debug mode, worm could give it commands to execute
 - sendmail ran the malicious code
- Password cracking with a dictionary of 432 words
 - accounts tested against words in a random order



"People pick bad passwords, and either forget, write down, or resent good ones."

Steven M. Bellovin



Effect of worm

- Formation of CERT
- \$10,000 fine, 3 year probation, and 400 hours of community service for Morris
- Heightened awareness of computer system vulnerabilities
- Something for security professionals to quote
 - not so much a problem now ☹️

<https://nvd.nist.gov/general/nvd-dashboard>



Example: Stuxnet (2010)

- Worm that propagates via USB sticks on Windows PCs
 - Actual target is a particular model of Siemens PLC
 - Used in many embedded industrial control systems
 - Exploits multiple vulnerabilities, including four “zero-day” exploits
 - Looks for Siemens SIMATIC WinCC/Step 7 controller software
 - If found, infects controller software (using previously unknown, unpatched vulnerability)
 - Read and alters bits of data in the controlled PLC's
- Unusually complex, costly to develop, lots of speculation
 - Eventually confirmed (anonymous US officials) as cyberattack on Iran's nuclear program



Example: CryptoLocker (2013)

- Propagates via email attachments
- Once activated:
 - Encrypts data files on the infected computer using strong RSA public-key cryptography
 - Sends a ransom note demanding the computer owner pay for the private decryption key via Bitcoin
 - Attempts to delete Windows Shadow Copy backups before encrypting files
- Offline backup or online pay-up...



Example: Heartbleed (2014)

- Vulnerability in OpenSSL cryptographic software library
- Bug is in the OpenSSL's implementation of the TLS/DTLS (transport layer security protocols) heartbeat extension
 - Improper input validation (due to a missing bounds check)
- Leads to the leak of memory contents from the server to the client and from the client to the server
 - In particular, leaks private keys


(<https://xkcd.com/424/>)



Ex: Meltdown, Spectre, Foreshadow (2018)



MELTDOWN



SPECTRE




FORESHADOW

Examples of *side-channel attacks*



Exploit CPU Performance Features

- Caching, speculative execution, out-of-order execution



- Meltdown: Exploits out-of-order execution
- Spectre: Exploits branch prediction



Meltdown in a Nutshell

1. Load from protected address into pointer
2. Use value to index into an array
3. Time reads from array to detect which entry was cached -> index of fast entry is value at protected memory

See slides from Foreshadow talk at USENIX Security 2018
(https://www.usenix.org/sites/default/files/conference/protected-files/security18_slides_bulck.pdf)



Spectre in a Nutshell

1. Run attacker code on same core as victim code
2. Train branch predictors to drive speculation down a particular branch in victim
3. Apply timing analysis to extract victim data



Buggy Code

- 85% of CERT Advisories describe problems that cannot be fixed with cryptography.
- Most of these are bugs in code
- But writing correct code is the oldest -- and probably the most difficult -- problem in computer science.
 - We're not going to solve it any time soon, possibly not ever.
- Reduction strategies
 - Structure – isolate security critical code; Safer languages
 - Reducing code complexity



Le Fin de Semester

FINALS BE LIKE:

**GOOD NIGHT STUDENT. GOOD WORK. SLEEP WELL.
I'LL MOST LIKELY KILL YOU IN THE MORNING.**

Good luck on all your exams!



From bug to vulnerability

- So you have a buggy user-level application
 - Why is this so bad?
- In general, compromising a process allows attacker to obtain privileges of that process for arbitrary activities
 - Bad for you, but not necessarily bad for the system
 - Compromising a process with root privilege (on Unix systems) provides a lot of power
 - Read/write any file
 - Read/write kernel memory though `/dev/kmem`
 - Attach to and trace any running process
 - Install kernel modules / change system configuration



Insufficient Domain Separation

- Authorization domains should be clearly separated
 - otherwise, less-privileged code can get more-privileged code to do bad things
- Unfortunately, this is often not the case
- Examples:
 - environment inheritance by setuid programs in UNIX
 - e.g., max file length or number of files open



Security policies are critical

- Most organizations have a stated policy about control of private information and access to resources
- These policies can help guide protocol implementation
 - and can help with political and “clueless user” problems
- If you can't say what's important, how am I supposed to protect it?
 - ... and why should I bother trying?



OS Security Mechanisms

- Access Controls
- FreeBSD Jails
- Flask Security Architecture (SELinux)



Access Control

- Common Assumption:
 - System knows identity of user (authentication)
 - Access requests pass through some gatekeeper (authorization)
- Implemented using Access Control Matrix
 - Access control list
 - Capability
- Two main types
 - Discretionary Access Control (DAC)
 - User sets access rights for objects they own
 - Mandatory Access Control (MAC)
 - System sets rights that users can't override



FreeBSD Jails

- Goal: isolation of processes to contain possible damage without lots of extra security management complexity
- Built on chroot concept
 - Give process (and all its children) separate view of file system tree (`chdir /tmp/limited_fs; chroot /tmp/limited_fs`)
 - Originally introduced for development
- Added new “jail” command
 - Each jail has own superuser
 - Privileges of superuser restricted to only affect things inside jail
 - Process in jail isolated from ones outside jail



Flask Security Architecture: Motivation

- No single definition of security suffices
- Need for many policies and even types of policies
- Computer security solutions must be flexible enough to support wide range of security policies
- This policy flexibility must be supported by the OS mechanisms



Defining Policy Flexibility

- Can't define through a list of known policies
- Defined in context of a state machine model
 - atomic operations to transition from one state to next
 - policy can interpose atomically on set of controlled operations
 - policy may use knowledge of portion of system state
- 3 Requirements of Policy Flexibility
 - Support fine-grained access controls on low-level objects
 - Propagate access rights according to security policy
 - Deal with changes in policy over time, including revoking previously granted permissions



Policy Changes

- Even simplest policies undergo changes
- Risk of enforcing obsolete policy
- Need for effective atomicity in policy changes
- Complicated by migrated permissions
 - access rights explicitly cached in data structures
 - access rights implicitly cached by operations in progress



Popular Mechanisms are Insufficient

- Capability-based systems
 - propagation of access rights
 - Hydra, KeyKOS, EROS: provide enhancements to limit propagation, but still lacking in support for policies
 - SCAP, ICAP, TMach: do not define mechanisms by which policy is queried to validate capabilities
- Interposition
 - mismatch between functional interface and security needs
 - does not support revocation of migrated permissions

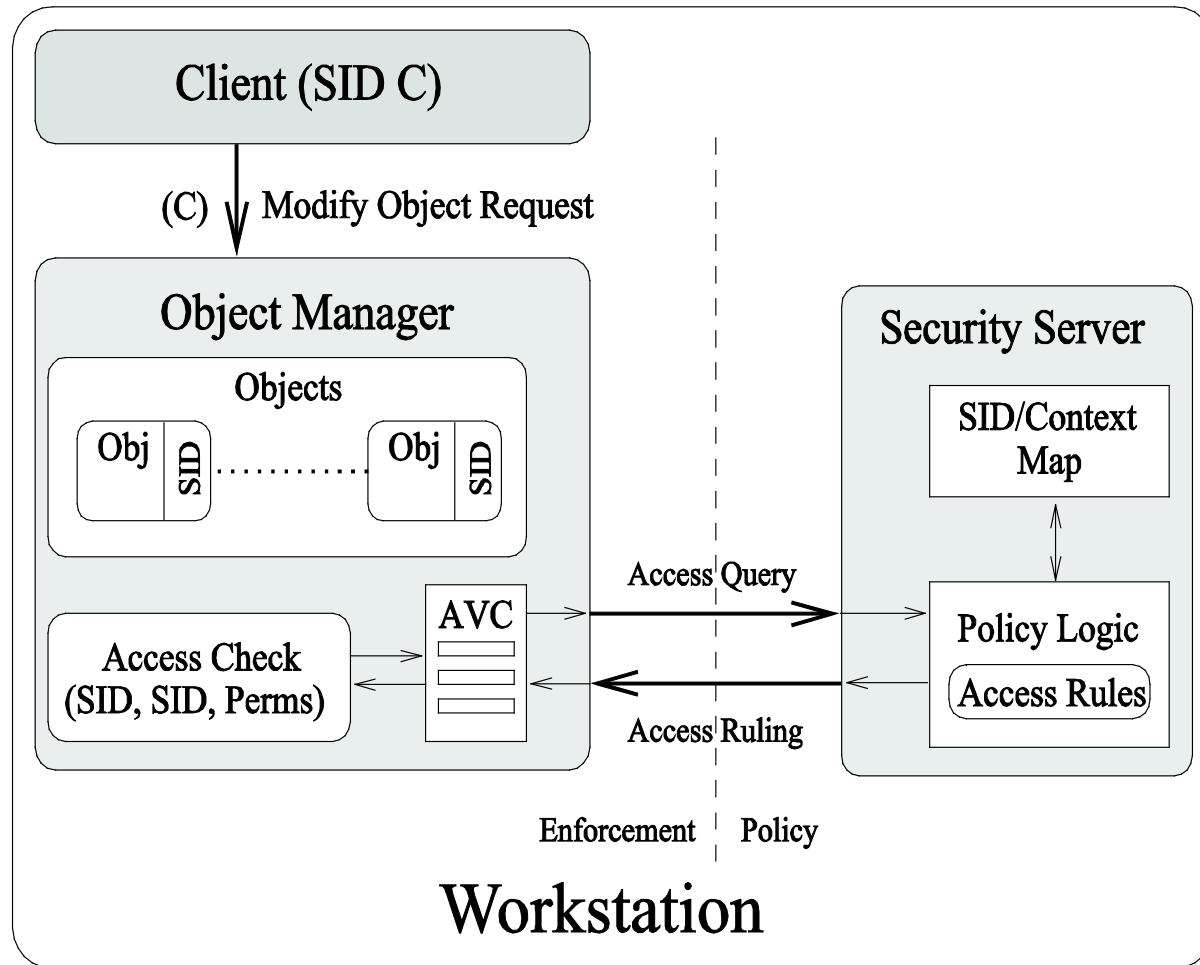


Flask Architecture

- Security server
 - provides labeling, access and polyinstantiation decisions
 - security contexts and security identifiers (SIDs)
- Object managers
 - bind labels to objects
 - enforce access decisions
 - direct clients to appropriate instances
- Access Vector Cache (AVC) library
 - coordinates access decisions
 - minimizes performance impact
- Underlying IPC mechanism
 - must provide identification of clients and servers



Architecture Diagram





Revocation Support

- Object manager atomicity requirements
 - Completeness
 - Promptness
- Protocol between security server and object managers
 - system-wide atomicity for policy changes
 - Security server notifies object manager AVC modules
 - AVC modules update cache state, invoke callbacks
 - Callbacks update migrated permissions
 - AVC modules notify security server of completion



Evaluation of Flexibility

- Support for policy changes
 - architecture provides support for system-wide atomicity
 - microkernel meets object manager atomicity requirements
 - other object managers lack support for migrated permissions
- Set of operations controlled by policy
 - fine-grained controls over all object services
- Set of operations that may be invoked by policy
 - object manager interfaces, AVC module interface
- System state available to policy
 - SID pairs sufficient for most policies (DTOS)
 - use prototype to research need for richer interface



Current Status

- SELinux – Security Enhanced Linux
 - Version of Linux created by the NSA and Secure Computing Corporation (SCC)
 - Incorporates University of Utah Flask security model
 - Supports mandatory access control to all objects
 - Separates Object Managers from Security Server
 - Supports various policy configurations
 - Often criticized for being too complicated
- TrustedBSD
 - Part of this system is a port of SELinux extensions to FreeBSD
 - TrustedDarwin is a port of TrustedBSD to the Darwin system
 - Some components of TrustedBSD have spilled over into OS X
 - not sure if this includes Flask implementation



seccomp

- Sandbox mechanism available in Linux
- Performs filtering on system calls
 - Filters are written in Berkeley Packet Filter (BPF) language
 - Attached by a process to itself via `prctl(PR_SET_SECCOMP, SECCOMP_MODE_FILTER, filter);`
 - Inherited by children
- Filter code can access system call number, arguments
 - Allows or denies the system call



Uses?